

Big Data Technologies

Agenda

- Hadoop 1.x vs 2.x
- Reading hadoop logs
- Job counters
- MR job configuration
- Uber job
- Input/Output format
- Cascaded MR jobs
- ~~MR execution with single reducer~~
- ~~Combiner~~
- ~~Partitioner~~
- ~~MR execution with multiple reducer~~
- ~~Map-only job~~
- ~~Reduce-only job~~
- ~~MR execution on YARN~~
- ~~Hadoop Performance~~
- ~~Hadoop streaming~~

Code Sharing

- <http://172.18.4.63:4200>

Hadoop 1.x vs 2.x

- Hadoop 1.x -- HDFS + MapReduce
 - HDFS: NameNode, DataNode, SecondaryNameNode
 - MapReduce: JobTracker, TaskTracker
 - JobTracker -- Keep track of progress of whole job (mapper(s) + reducer(s)).

- One JobTracker in the cluster (Master).
- TaskTracker -- Keep track of progress of individual tasks i.e. mapper/reducer.
 - Multiple TaskTrackers in the cluster (Workers).
- MR job/processing is equally distributed across the cluster/TaskTrackers by the Master/JobTracker.
- Hadoop 2.x -- HDFS + MapReduce on YARN
 - HDFS: NameNode, DataNode, SecondaryNameNode (and/or StandbyNameNode, JournalNode - HA)
 - YARN: ResourceManager, NodeManager
 - ResourceManager -- Keep track of all computing resources in the cluster.
 - One ResourceManager in the cluster (Master) -- Multiple ResourceManagers can be configured for HA.
 - NodeManager -- Keep track of execution of tasks on individual worker node.
 - Multiple NodeManagers in the cluster (Workers).
 - MR job/processing is distributed across the cluster depending on availability of the resources.
 - Can use one of the following scheduler:
 - FIFO: Executes only one job at a time. Jobs executed in order of arrival.
 - Capacity: Can execute multiple jobs in parallel depending on availability of the computing resources.
 - Fair Scheduler: Can execute multiple jobs in parallel. If few resources are available, then schedule small jobs there.

Ncdc Job execution

- terminal> start-dfs.sh && start-yarn.sh
- terminal> mapred --daemon start historyserver
- browser> http://localhost:19888
- terminal> hadoop fs -mkdir -p /user/nilesh/ncdc/input
- terminal> hadoop fs -put * /user/nilesh/ncdc/input #command to be executed from data/ncdc directory
- terminal> hadoop jar mr3-ncdc-0.0.1-SNAPSHOT.jar /user/nilesh/ncdc/input /user/nilesh/ncdc/output1
- terminal> hadoop fs -ls /user/nilesh/ncdc/output1
- terminal> hadoop fs -head /user/nilesh/ncdc/output1/part-r-00000

Read hadoop logs

- Method 1: Access the worker nodes (over ssh/direct terminal)
 - \$HADOOP_HOME/logs

- Hadoop Daemon logs
- --> userlogs --> Hadoop MR application logs
 - --> application directory
 - --> one directory for one task/container -- MrAppMaster, Mapper(s), Reducer(s).
- Method 2: Job historyserver
 - http://localhost:19888
 - Tools --> Local logs
 - Hadoop Daemon logs
 - --> userlogs --> Hadoop MR application logs
 - --> application directory
 - --> one directory for one task/container -- MrAppMaster, Mapper(s), Reducer(s).

Job Counters

- Hadoop maintains various system/pre-defined counters while executing MR jobs.
 - File System Counters
 - Job Counters -- Number of Mappers/Reducers, ...
 - Map-Reduce Framework -- Number of Mapper Input records, ...
 - Shuffle Errors
 - File Input/Output Format Counters
- To add custom job counter, create an enum.

```
enum NcdcJobCounter {  
    VALID_READING, INVALID_READING, INVALID_RECORD  
}
```

- To increment the counter

```
Counter cnt = context.getCounter(NcdcJobCounter.VALID_READING);  
cnt.increment(1);
```

- See output after executing job on terminal or in job historyserver.

MR job configuration

- MR job configuration can be given while executing MR on command line.
 - `hadoop jar app.jar -fs hdfs://master:9000 -jt master:8032 -D dfs.replication=2 arg1 arg2 ...`
- However it is common practice to put all required settings in a XML file and specify it using `-conf` generic option.
 - `hadoop jar app.jar -conf job-config.xml arg1 arg2 ...`
- Typically we can maintain different config files for executing same program in different environments.
 - `local-config.xml` -- execute on single machine (no HDFS, no YARN).
 - No need to start HDFS or YARN.
 - `hadoop jar app.jar -conf job-config.xml /home/$USER/bigdata/data/ncdc/tmp/output1`
 - `ls /tmp/output1`
 - `pseudo-config.xml` -- execute on single node hadoop cluster.
 - `cluster-config.xml` -- execute on multi node hadoop cluster.
- Usually config includes performance tuning settings for the job e.g. mapper output buffer size, number of reducers, jvm memory size, etc.

Uber job

- By default when a MR job executes on cluster, at least three containers/JVM processes are created i.e. MrAppMaster, Mapper, and Reducer.
- If job processing/computing is quite "small", we can execute all these tasks (MrAppMaster, Mappers, Reducers) in single JVM process/container, so that there is no IPC needed and execution is much faster. This will also save computing resources. Such execution is called as "uber" job.
- The small job can be defined in the config. The default config

```
<property>
  <name>mapreduce.job.ubertask.enable</name>
  <value>>false</value>
</property>
<property>
  <name>mapreduce.job.ubertask.maxmaps</name>
  <value>9</value>
</property>
```

```
<property>
  <name>mapreduce.job.ubertask.maxreduces</name>
  <value>1</value>
</property>
```

- To enable uberization of task, we should modify the above config as per our requirement. Also programmer should keep computing resources of cluster in mind and configure them too if needed.
- Refer uber-config.xml.

```
<property>
  <name>mapreduce.job.ubertask.enable</name>
  <value>true</value>
</property>
```

- terminal> `hadoop jar mr3-ncdc-0.0.1-SNAPSHOT.jar -conf ~/dbda/bigdata/day06/mr-config/uber-config.xml /user/nilesh/ncdc/input /user/nilesh/ncdc/output5`

Input/Output format

- InputFormat
 - Used to read input data record by record from HDFS.
 - To read each record it internally uses "RecordReader" which is a nested class of InputFormat.
 - TextInputFormat and its RecordReader -- each record is one line (ended by '\n' or '\r').
 - Key=LongWritable -- line offset
 - Value=Text -- line
 - NLineInputFormat and its RecordReader -- each split is "n" lines.
 - Below example: one record/split = 3 lines.

```
name: Nilesh Ghule
email: nilesh@sunbeaminfo.com
```

```
mobile: 9527331338
name: Nitin Kudale
email: nitin@sunbeaminfo.com
mobile: 9881208115
name: Prashat Lad
email: prashant@sunbeaminfo.com
mobile: 9881208114
```

- `KeyValueTextInputFormat` and its `RecordReader` -- each record is one line with key and value separated by tab.
 - Key=Text before tab(separator)
 - Value=Text after tab(separator)

```
1  -74.49494451294697
2  -75.16562866684718
3  -44.24810668422786
4   9.304906241972771
5   67.9370272936534
6  126.61202558635395
7  160.3485770685968
8  137.20253584155134
9   91.14684224205146
10  36.454816285998014
11 -15.835389526356655
12 -56.066257566962435
```

- `CombineTextInputFormat` & its `RecordReader` -- treat multiple files as a single split (if size is smaller than block size).
 - e.g. If 20 files with total data size < 100mb, will be treated as single input split. So will create single mapper and execution will be faster.
- `FixedLengthInputFormat` & its `RecordReader` -- treat "n" bytes as one record.
 - Useful for text/binary file where each record is fixed number of bytes.
- `SequenceFileInputFormat` & its `RecordReader` -- used to read sequence files.
 - These files have a definite structure. They are binary and can be compressed.
- `DBInputFormat` & its `RecordReader` -- used to read records from RDBMS.

- **OutputFormat**
 - Used to write output data record by record into HDFS.
 - To write each record it internally uses "RecordWriter" which is a nested class of OutputFormat.
 - **TextOutputFormat**
 - each record: key "tab" value "\n"
 - **SequenceFileOutputFormat**
 - To write records in binary and compressed (optional) format
 - **DBOutputFormat**
 - To write output records into RDBMS.

Max Avg Temperature NCDC

- **Input: Monthwise average Temperature. (Output of Avg Temperature job)**

```
1 -74.49494451294697
2 -75.16562866684718
3 -44.24810668422786
4 9.304906241972771
5 67.9370272936534
6 126.61202558635395
7 160.3485770685968
8 137.20253584155134
9 91.14684224205146
10 36.454816285998014
11 -15.835389526356655
12 -56.066257566962435
```

- **Output: Maximum Temperature and its Month**

```
7 160.3485770685968
```

- Mapper:

- Code/Logic

```
// if using KeyValueTextInputFormat -- Mapper input will be Text key, Text value
class MaxTemperatureMapper extends Mapper<Text,Text,NullWritable,Text> {
    @Override
    public void map(Text key, Text value, Mapper.Context context) {
        int month = Integer.parseInt(key.toString());
        double temperature = Double.parseDouble(value.toString());
        String monthTemperature = month + "," + temperature;
        context.write(NullWritable.get(),
            new Text(monthTemperature));
    }
}
```

- Mapper output

```
x    1, -74.49494451294697
x    2, -75.16562866684718
x    3, -44.24810668422786
x    4, 9.304906241972771
x    5, 67.9370272936534
x    6, 126.61202558635395
x    7, 160.3485770685968
x    8, 137.20253584155134
x    9, 91.14684224205146
x   10, 36.454816285998014
x   11, -15.835389526356655
x   12, -56.066257566962435
```

- Reducer:

- Code/Logic

```
class MaxTemperatureReducer extends Reducer<NullWritable,Text,IntWritable,DoubleWritable> {  
    @Override  
    public void reduce(NullWritable key, Iterable<Text> values, Reducer.Context context) {  
        double maxTemperature = -Double.MIN_VALUE;  
        int maxMonth = 0;  
        for(Text monthTemperatureWr: values) {  
            String monthTemperature = monthTemperatureWr.toString();  
            String[] parts = monthTemperature.split(",");  
            int month = Integer.parseInt(parts[0]);  
            double temperature = Double.parseDouble(parts[1]);  
            if(temperature > maxTemperature) {  
                maxTemperature = temperature;  
                maxMonth = month;  
            }  
        }  
        context.write(new IntWritable(maxMonth), new DoubleWritable(maxTemperature));  
    }  
}
```

- Output

```
7    160.3485770685968
```

Cascaded MR Job

- Input --> Job1 --> Aux Output --> Job2 --> Output
- Driver code needs to be modified creating and submitting second job.
- Second job must be submitted only if first job is successful.
- Typically, input directory of Second job is same as output directory of First job.

Hadoop Writables

- NullWritable -- for keeping key or value null.
- To send multiple fields in "key" or "value", we can do one of the following.
 - option 1: Use "Text" that keeps multiple values concatenated and separated using some symbol (e.g. ",").
 - option 2: To send multiple values of the same type, use ArrayWritable.
 - option 3: To send two value of different type, use MapWritable.
 - option 4: To send multiple value as "value", write custom writable class.

```
class CustomWritable extends Writable {  
    private int month;  
    private double temperature;  
    // constructors  
    // getters/setters  
    public void write(DataOutput dout) {  
        dout.write(month);  
        dout.write(temperature);  
    }  
    public void readFields(DataInput din) {  
        this.month = din.readInt();  
        this.temperature = din.readDouble();  
    }  
}
```

- option 5: To send multiple value as "key", write custom writable class.

```
class CustomWritable extends WritableComparable<CustomWritable> {  
    private int month;  
    private double temperature;  
    // constructors  
    // getters/setters  
    public void write(DataOutput dout) {
```

```
        dout.write(month);
        dout.write(temperature);
    }
    public void readFields(DataInput din) {
        this.month = din.readInt();
        this.temperature = din.readDouble();
    }
    public int compareTo(CustomWritable other) {
        // ...
    }
    // equals() and hashCode()
}
```

Assignment

- NCDC Data --> Job1 --> Yearly Avg Temperature --> Job2 --> Max & Min Temperature & Year.

Java Reflection Tutorial

- `ClassName.class`
- https://youtu.be/lAoNJ_7LD44