

Big Data Technologies

Agenda

- Hive

ORC format

- Optimized Row Columnar format
- Designed by Hive team for efficient execution of hive queries
- Hadoop has ORCInputFormat and ORCOutputFormat.
- Hive queries when work on ORC tables, they internally use ORCInputFormat to read data into mapper and use ORCOutputFormat to write reducer output data into HDFS.

Vectorization

- To perform math operations on primitive types, hive uses vectorization.
- Typically Hive data is stored in ORC format, that is further divided into the blocks and arranged in columns.
- This helps speeding up math operations on primitive types (block by block).

```
set hive.vectorized.execution.enabled=true;
```

Hive Architecture

- Refer slides

Hive Compaction

- Minor compaction: Smaller changes files (delta) are merged together to reduce the number of delta files. This will save space in HDFS and also speed-up further queries. In minor compaction, data in main file is not modified (e.g. older records are still present there). Minor compaction process takes

relatively less time.

- Major compaction: All delta files and main data files are merged to create a new data file. All older versions (delta files and older main file) are deleted. This will significantly improve speed to further queries. Major compaction needs more time.

```
ALTER TABLE ncdc_orc COMPACT 'major';
```

Hive scripts

- Set of hive commands.
- Example 1: Compaction process

```
-- compact.hql  
ALTER TABLE ncdc_orc COMPACT 'major';
```

- Example 2: Rebuild correlation table

```
-- rebuild-correlation.hql  
SET mapreduce.reduce.memory.mb = 5120;  
SET mapreduce.reduce.java.opts = -Xmx4096m;  
  
INSERT INTO movies_orc  
SELECT * FROM movies_staging WHERE condition_to_get_new_records_only;  
  
INSERT INTO ratings_orc  
SELECT * FROM ratings_staging WHERE condition_to_get_new_records_only;  
  
ALTER MATERIALIZED VIEW mv_user_movies REBUILD;  
  
DROP TABLE IF EXISTS movies_corr;
```

```
CREATE TABLE movies_corr AS
SELECT m1, m2, COUNT(rt1) cnt, CORR(rt1,rt2) cor
FROM mv_user_movies
GROUP BY m1, m2;
```

- To run the script:
 - beeline> !run /path/of/script.hql
 - OR
 - terminal> beeline -u jdbc:hive2://localhost:10000/dbda -n \$USER -f /path/of/script.hql

Managed Table vs External Table

- Managed Table
 - CREATE TABLE tablename ...
 - Created in HDFS "warehouse" directory and metadata is stored in metastore.
 - To load the data in the table -- LOAD DATA or INSERT.
 - DROP TABLE tablename ... -- drop the data from HDFS and metadata from metastore.
- External Table
 - CREATE EXTERNAL TABLE tablename ... LOCATION '/path/of/data/directory';
 - Data is already present in HDFS (not necessarily in "warehouse") and metadata is stored in metastore.
 - The data is already present -- LOCATION '...' -- given file table creation. To load the additional data in the table -- LOAD DATA or INSERT.
 - DROP TABLE tablename ... -- drop metadata from metastore. Data from HDFS is not deleted.

```
CREATE EXTERNAL TABLE emp_staging(
empno INT,
ename STRING,
job STRING,
mgr INT,
hire STRING,
sal DOUBLE,
comm DOUBLE,
deptno INT
```

```
)  
ROW FORMAT DELIMITED  
FIELDS TERMINATED BY ','  
STORED AS TEXTFILE  
LOCATION '/user/nilesh/emp/input';  
  
SELECT * FROM emp_staging  
LIMIT 10;
```

Hive Functions

- Hive function types
 - Single row functions/Scalar functions
 - "n" rows --> Function --> "n" rows
 - Multi row functions/Group functions/Aggregated functions
 - "n" rows --> Function --> "m" rows ($m < n$)
 - Table valued functions
 - "n" rows --> Function --> "m" rows ($m > n$)
- Hive user-defined functions
 - UDF -- User Defined Function
 - UDAF -- User Defined Aggregate Function
 - UDTF -- User Defined Table Function
 - Implemented in Java.

```
DESCRIBE movies_orc;  
  
SELECT * FROM movies_orc LIMIT 5;  
  
SELECT id, title, SPLIT(genres, '\\|') FROM movies_orc  
LIMIT 5;  
  
CREATE TABLE movies(id INT, title STRING, genres ARRAY<STRING>)
```

```

STORED AS ORC;

INSERT INTO movies
SELECT id, title, SPLIT(genres, '\\|') FROM movies_orc;

SELECT * FROM movies
LIMIT 5;

SELECT COUNT(id) FROM movies
WHERE ARRAY_CONTAINS(genres, 'Romance');

```

- If input movie:

4	Waiting to Exhale (1995)	["Comedy", "Drama", "Romance"]
---	--------------------------	--------------------------------

- Output records

4	Waiting to Exhale (1995)	"Comedy"
4	Waiting to Exhale (1995)	"Drama"
4	Waiting to Exhale (1995)	"Romance"

```

SELECT id, title, EXPLODE(genres) FROM movies;
-- works in spark (not in hive)

SELECT id, title, genre FROM movies
LATERAL VIEW EXPLODE(genres) v_genres AS genre
LIMIT 20;

SELECT genre, COUNT(id) FROM movies

```

```
LATERAL VIEW EXPLODE(genres) v_genres AS genre
GROUP BY genre;
```

```
SELECT * FROM ratings_orc
LIMIT 5;

SELECT userid,movieid,rating,FROM_UNIXTIME(rtime) FROM ratings_orc
LIMIT 5;

CREATE TABLE ratings(
userid INT,
movieid INT,
rating DOUBLE,
rtime TIMESTAMP
)
STORED AS ORC;

INSERT INTO ratings
SELECT userid,movieid,rating,FROM_UNIXTIME(rtime) FROM ratings_orc;

SELECT * FROM ratings
LIMIT 5;

SELECT YEAR(rtime) yr, COUNT(rating) FROM ratings
GROUP BY YEAR(rtime);
```

Partitioning

Static partitioning

```
CREATE TABLE emp_part_dept(  
  empno INT,  
  ename STRING,  
  job STRING,  
  mgr INT,  
  hire STRING,  
  sal DOUBLE,  
  comm DOUBLE  
)  
PARTITIONED BY (deptno INT)  
ROW FORMAT DELIMITED  
FIELDS TERMINATED BY ','  
STORED AS TEXTFILE;  
  
DESCRIBE emp_part_dept;  
  
LOAD DATA LOCAL  
INPATH '/home/nilesh/sep22/dbda/bigdata/data/emp10.csv'  
INTO TABLE emp_part_dept  
PARTITION (deptno=10);  
  
LOAD DATA LOCAL  
INPATH '/home/nilesh/sep22/dbda/bigdata/data/emp20.csv'  
INTO TABLE emp_part_dept  
PARTITION (deptno=20);  
  
LOAD DATA LOCAL  
INPATH '/home/nilesh/sep22/dbda/bigdata/data/emp30.csv'  
INTO TABLE emp_part_dept  
PARTITION (deptno=30);  
  
EXPLAIN  
SELECT SUM(sal) FROM emp_staging  
WHERE deptno=20;
```

```
EXPLAIN
SELECT SUM(sal) FROM emp_part_dept
WHERE deptno=20;

DROP TABLE emp_part_dept;
```

Dynamic partitioning

```
CREATE TABLE emp_part_dept(
empno INT,
ename STRING,
job STRING,
mgr INT,
hire STRING,
sal DOUBLE,
comm DOUBLE
)
PARTITIONED BY (deptno INT)
ROW FORMAT DELIMITED
FIELDS TERMINATED BY ','
STORED AS TEXTFILE;

DESCRIBE emp_part_dept;

-- load emp.csv into a staging table -- emp_staging

INSERT INTO emp_part_dept PARTITION(deptno)
SELECT empno,ename,job,mgr,hire,sal,comm,deptno FROM emp_staging;
```

```
CREATE TABLE emp_part_dept_job(
empno INT,
```



```
ename STRING,  
mgr INT,  
hire STRING,  
sal DOUBLE,  
comm DOUBLE  
)  
PARTITIONED BY (deptno INT, job STRING)  
ROW FORMAT DELIMITED  
FIELDS TERMINATED BY ','  
STORED AS TEXTFILE;  
  
DESCRIBE emp_part_dept_job;  
  
INSERT INTO emp_part_dept_job PARTITION(deptno,job)  
SELECT empno,ename,mgr,hire,sal,comm,deptno,job FROM emp_staging;
```

Assignment

- External tables
 - Upload contacts.csv into HDFS.
 - Table contacts1: id INT, name STRING, emails STRING, addr STRING, phone STRING
 - SELECT * FROM contacts1 WHERE emails LIKE '%nilesh@sunbeaminfo.com%';
 - Table contacts2: id INT, name STRING, emails ARRAY<STRING>, addr STRUCT<...>, phone MAP<STRING,STRING>
 - SELECT * FROM contacts1 WHERE ARRAY_CONCATS(emails, 'nilesh@sunbeaminfo.com');
 - Table contacts3: line STRING
 - DROP TABLE contacts1;
- EXPLODE()
 - Load Hadoop license file into a Managed table line STRING.
 - Count number of occurrences for each word.
 - Count number of occurrences for each word other than stop words -- WHERE word NOT IN ('a', 'an', 'the', 'is', 'you', 'i', ...).
 - Find top 20 most frequent words.
- Dynamic Partitioning
 - Create emp table partitioned by job.

- Create another emp table partitioned by job and deptno.