



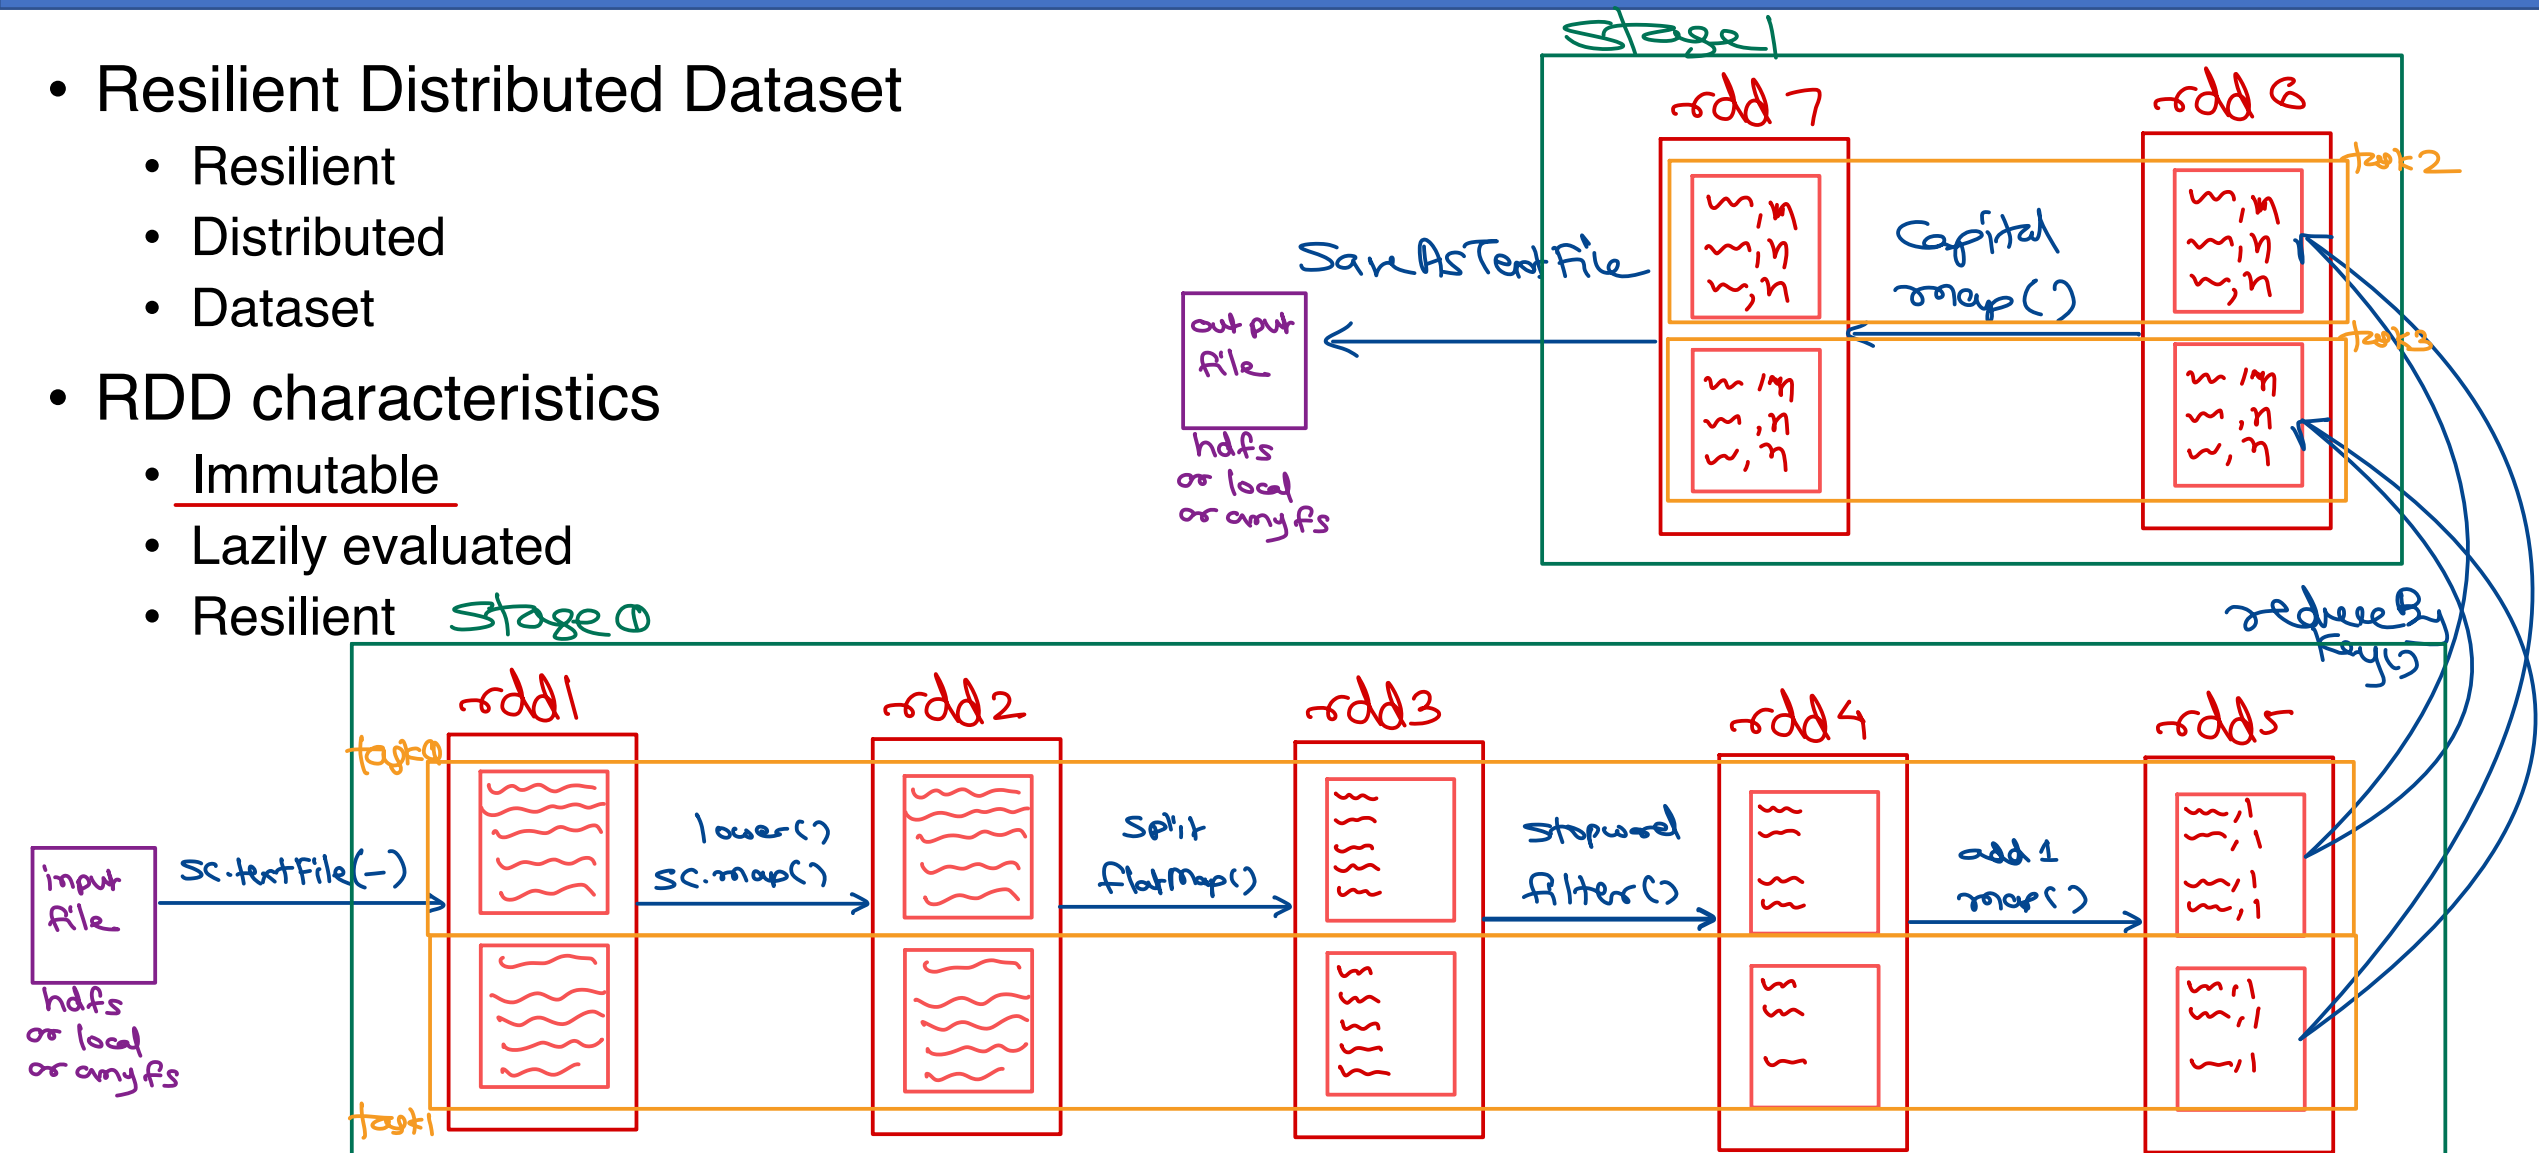
Apache Spark

Sunbeam Infotech



Spark RDD

- Resilient Distributed Dataset
 - Resilient
 - Distributed
 - Dataset
- RDD characteristics
 - Immutable
 - Lazily evaluated
 - Resilient

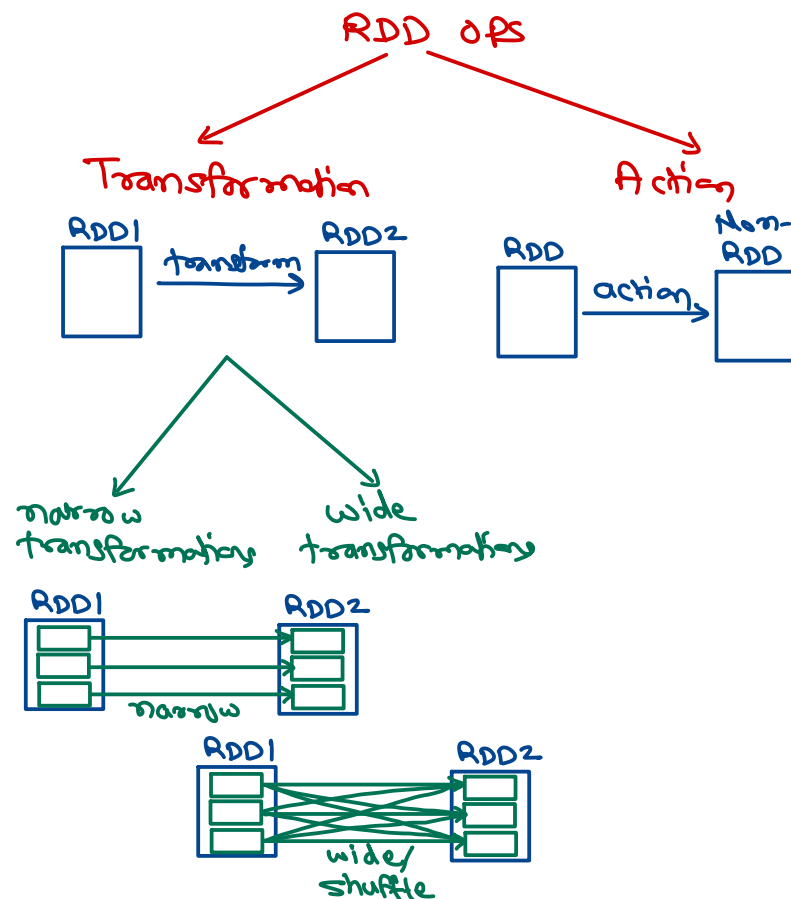


Spark RDD – creation

- sc.parallelize(collection, partitions)
 - convert collection into rdd with given partitions
- sc.textFile(path)
 - hdfs or local or s3 file or directory
- sc.wholeTextFiles(path)
 - hdfs or local or s3 directory, one file = one record
- sc.binaryRecords(path, recLen)
 - hdfs or local or s3 file or directory, recLen bytes = one record
- sc.wholeBinaryFiles(path)
 - hdfs or local or s3 directory, one file = one record
- sc.hadoopFile(path, inputFormat, ...)
 - hdfs file or directory, number of partitions = number of input splits

≈ number of hdfs blocks




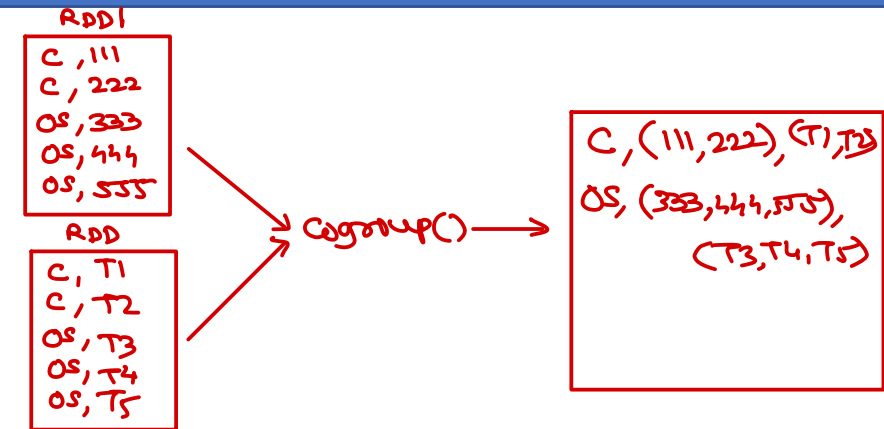
RDD Operations



RDD Operations

- Transformations: Returns RDD (*narrow + wide*)

- distinct()
- filter()
- map()
- flatMap()
- sortBy(keyValue, ascending, numOfPartitions) - *shuffle*
- sortByKey() - *shuffle*
- pipe() - *similar to hadoop streaming (execute program in other lang).*
- keyBy() → *add to KV pair RDD*
- countByKey() → *action*
- mapValues()
- groupByKey() → *shuffle*
- aggregateByKey() → *shuffle*
- cogroup()
- zip()
- join() – joins two RDDs by key.
- reduceByKey()

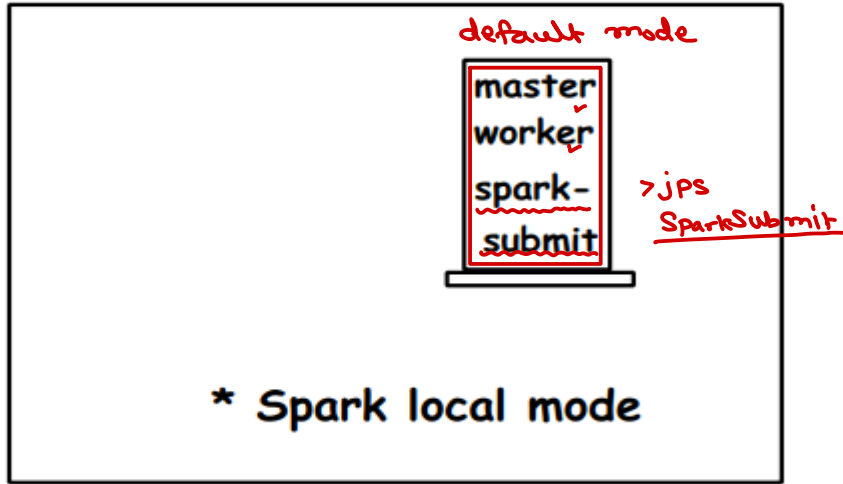
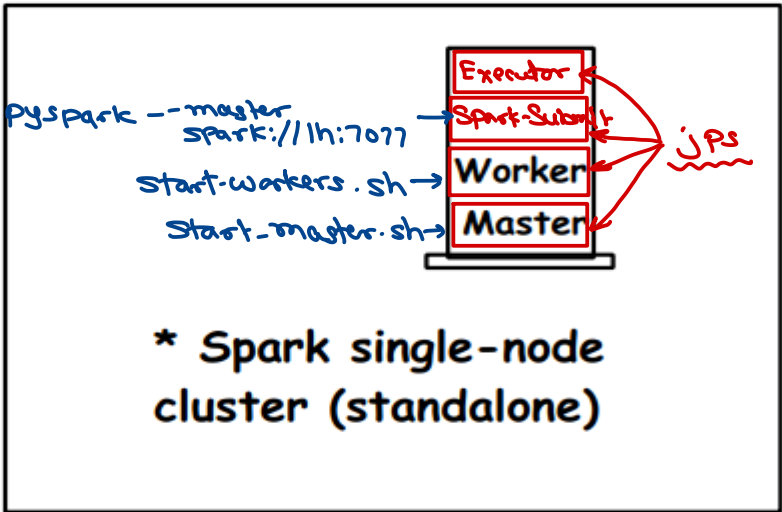
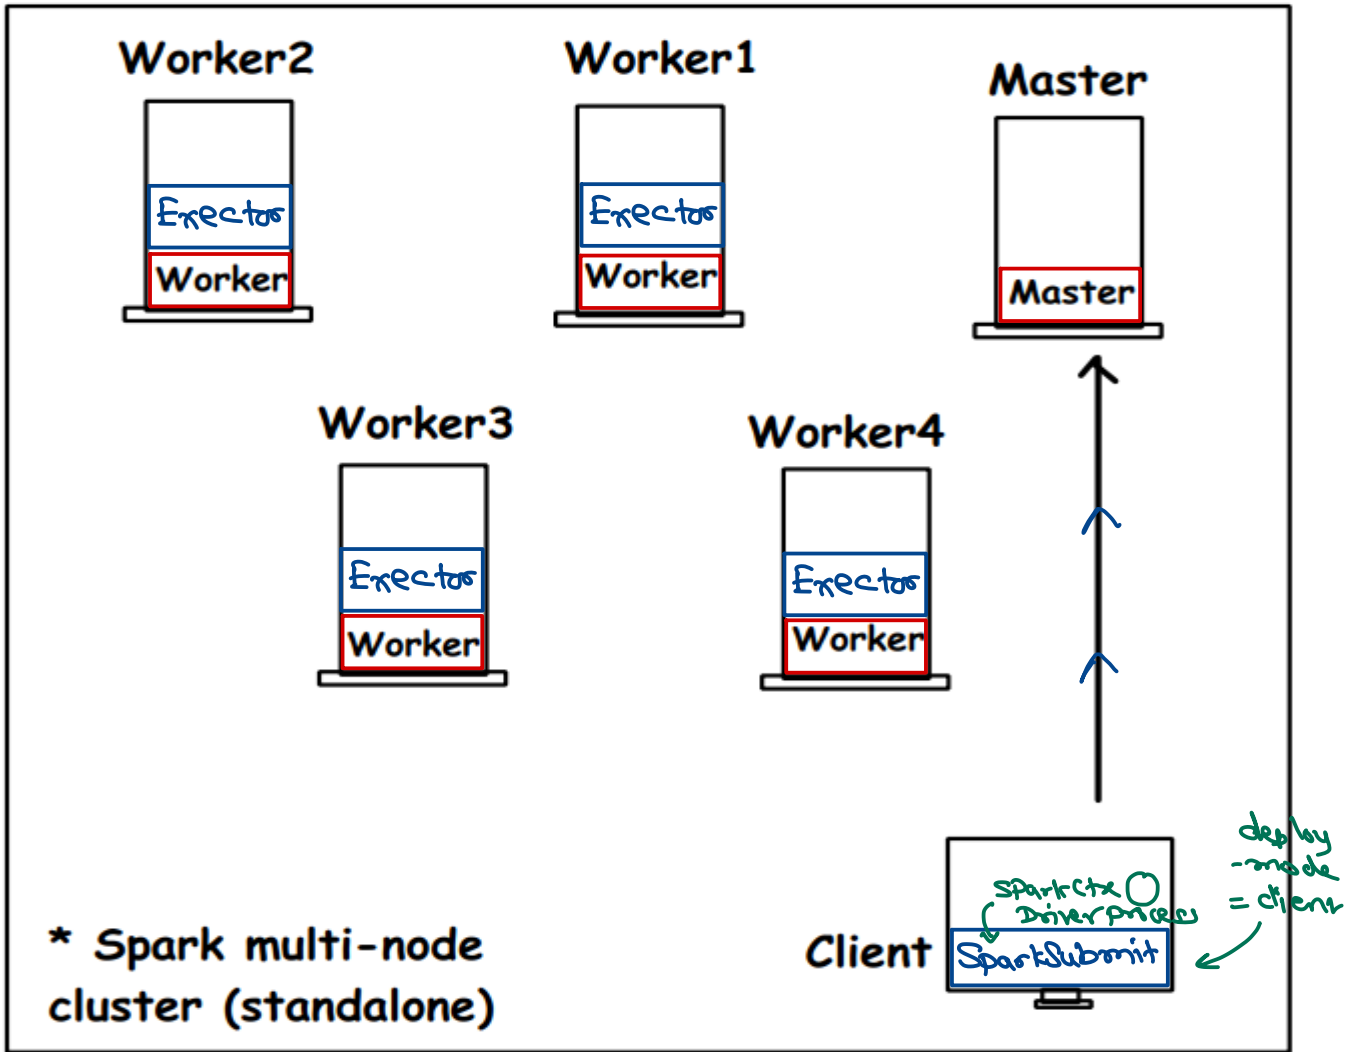


RDD Operations

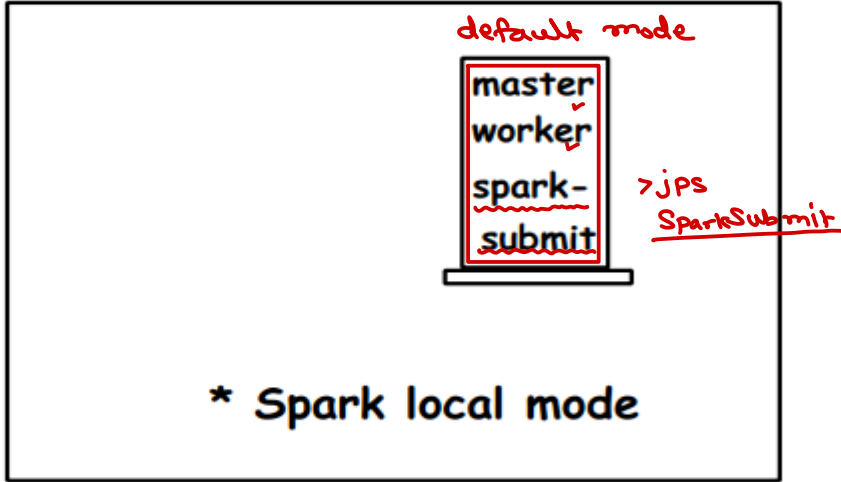
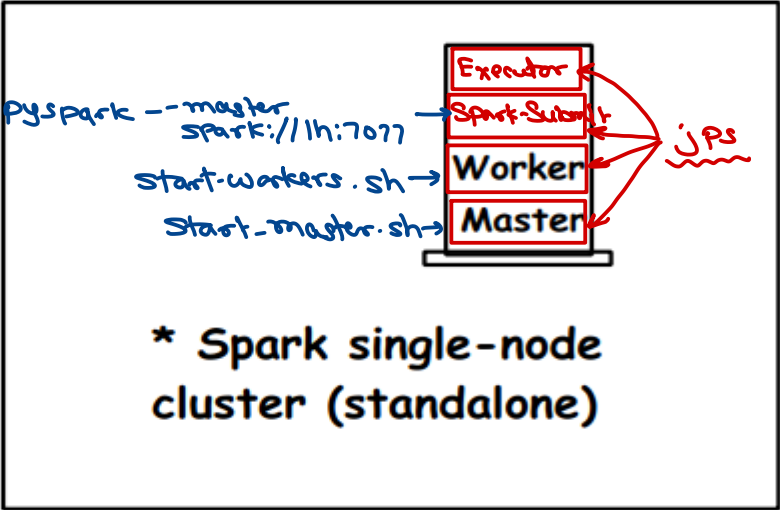
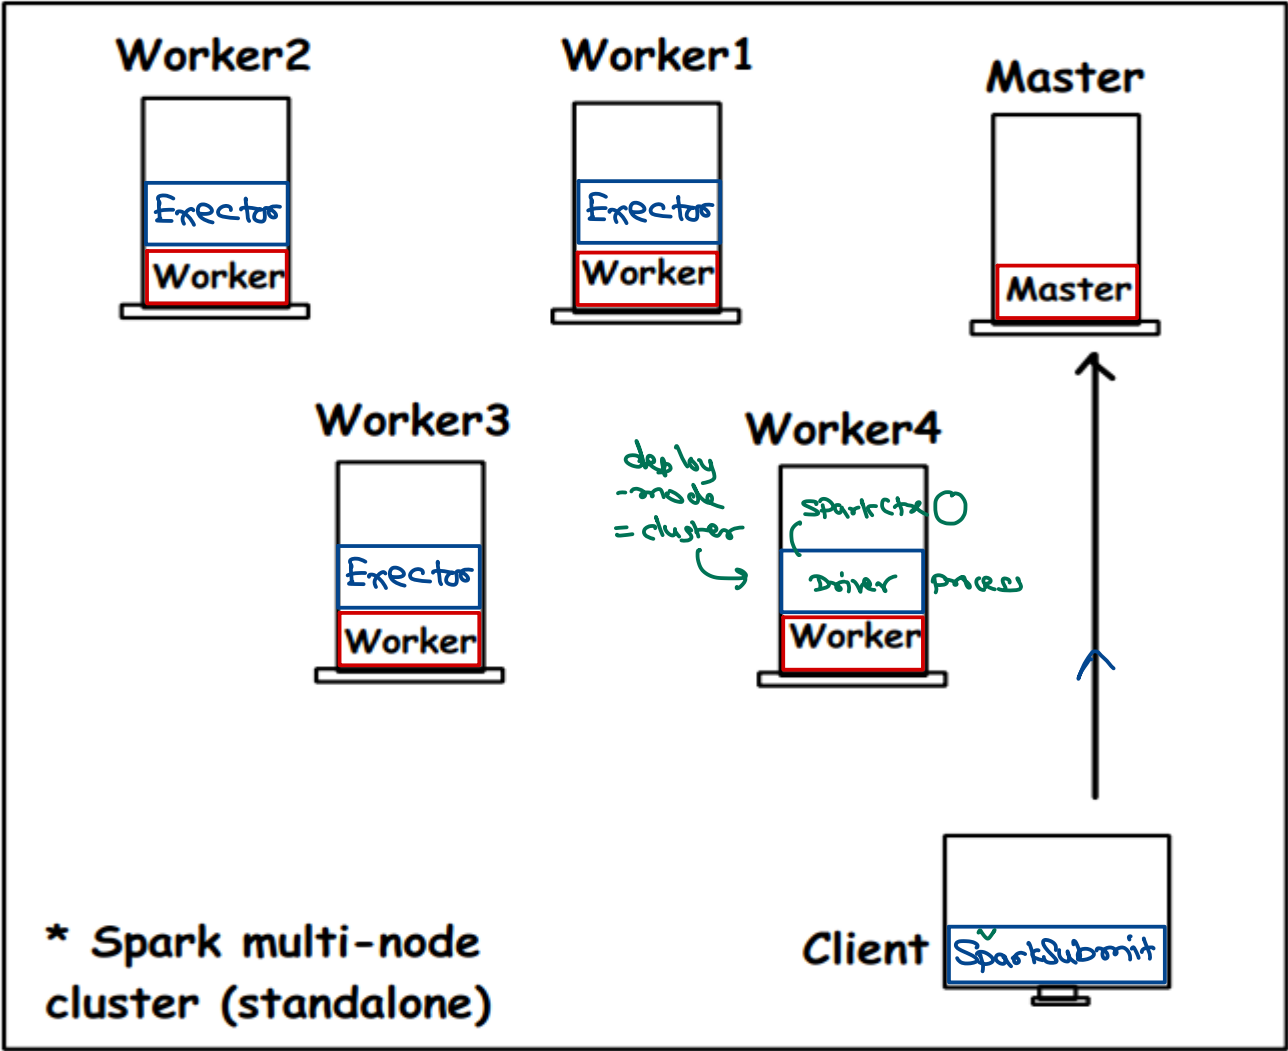
- Transformations
 - Narrow transformations
 - A partition of new RDD is computed from single partition of source RDD
 - Wide transformations
 - Partition of new RDD is computed from multiple partitions of source RDD
 - These transformations cause shuffling data across partitions.
- Actions: Returns non-RDD
 - count()
 - countApprox()
 - reduce()
 - countByValue()
 - first()
 - max()
 - min()
 - collect()
 - take() – collect n elements
 - takeOrdered()
 - top()
 - saveAsTextFile()
 - saveAsObjectFile() – *seq file*
 - lookup() – lookup by key



Spark Installation Modes



Spark Installation Modes



Spark Single Node Cluster

- Download & extract spark-x.y.z-bin-hadoop2.7.tgz.

- In ~/.bashrc

- export SPARK_HOME=\$HOME/spark-x.y.z-bin-hadoop2.7
- export PATH=\$SPARK_HOME/bin:\$SPARK_HOME/sbin:\$PATH

- In \$SPARK_HOME/conf/spark-env.sh

- export SPARK_MASTER_HOST=localhost
- export SPARK_LOCAL_IP=localhost

\$SPARK_HOME/conf/workers
localhost

- In \$SPARK_HOME/conf/spark-defaults.conf

- spark.master spark://localhost:7077

- Start master & slaves.

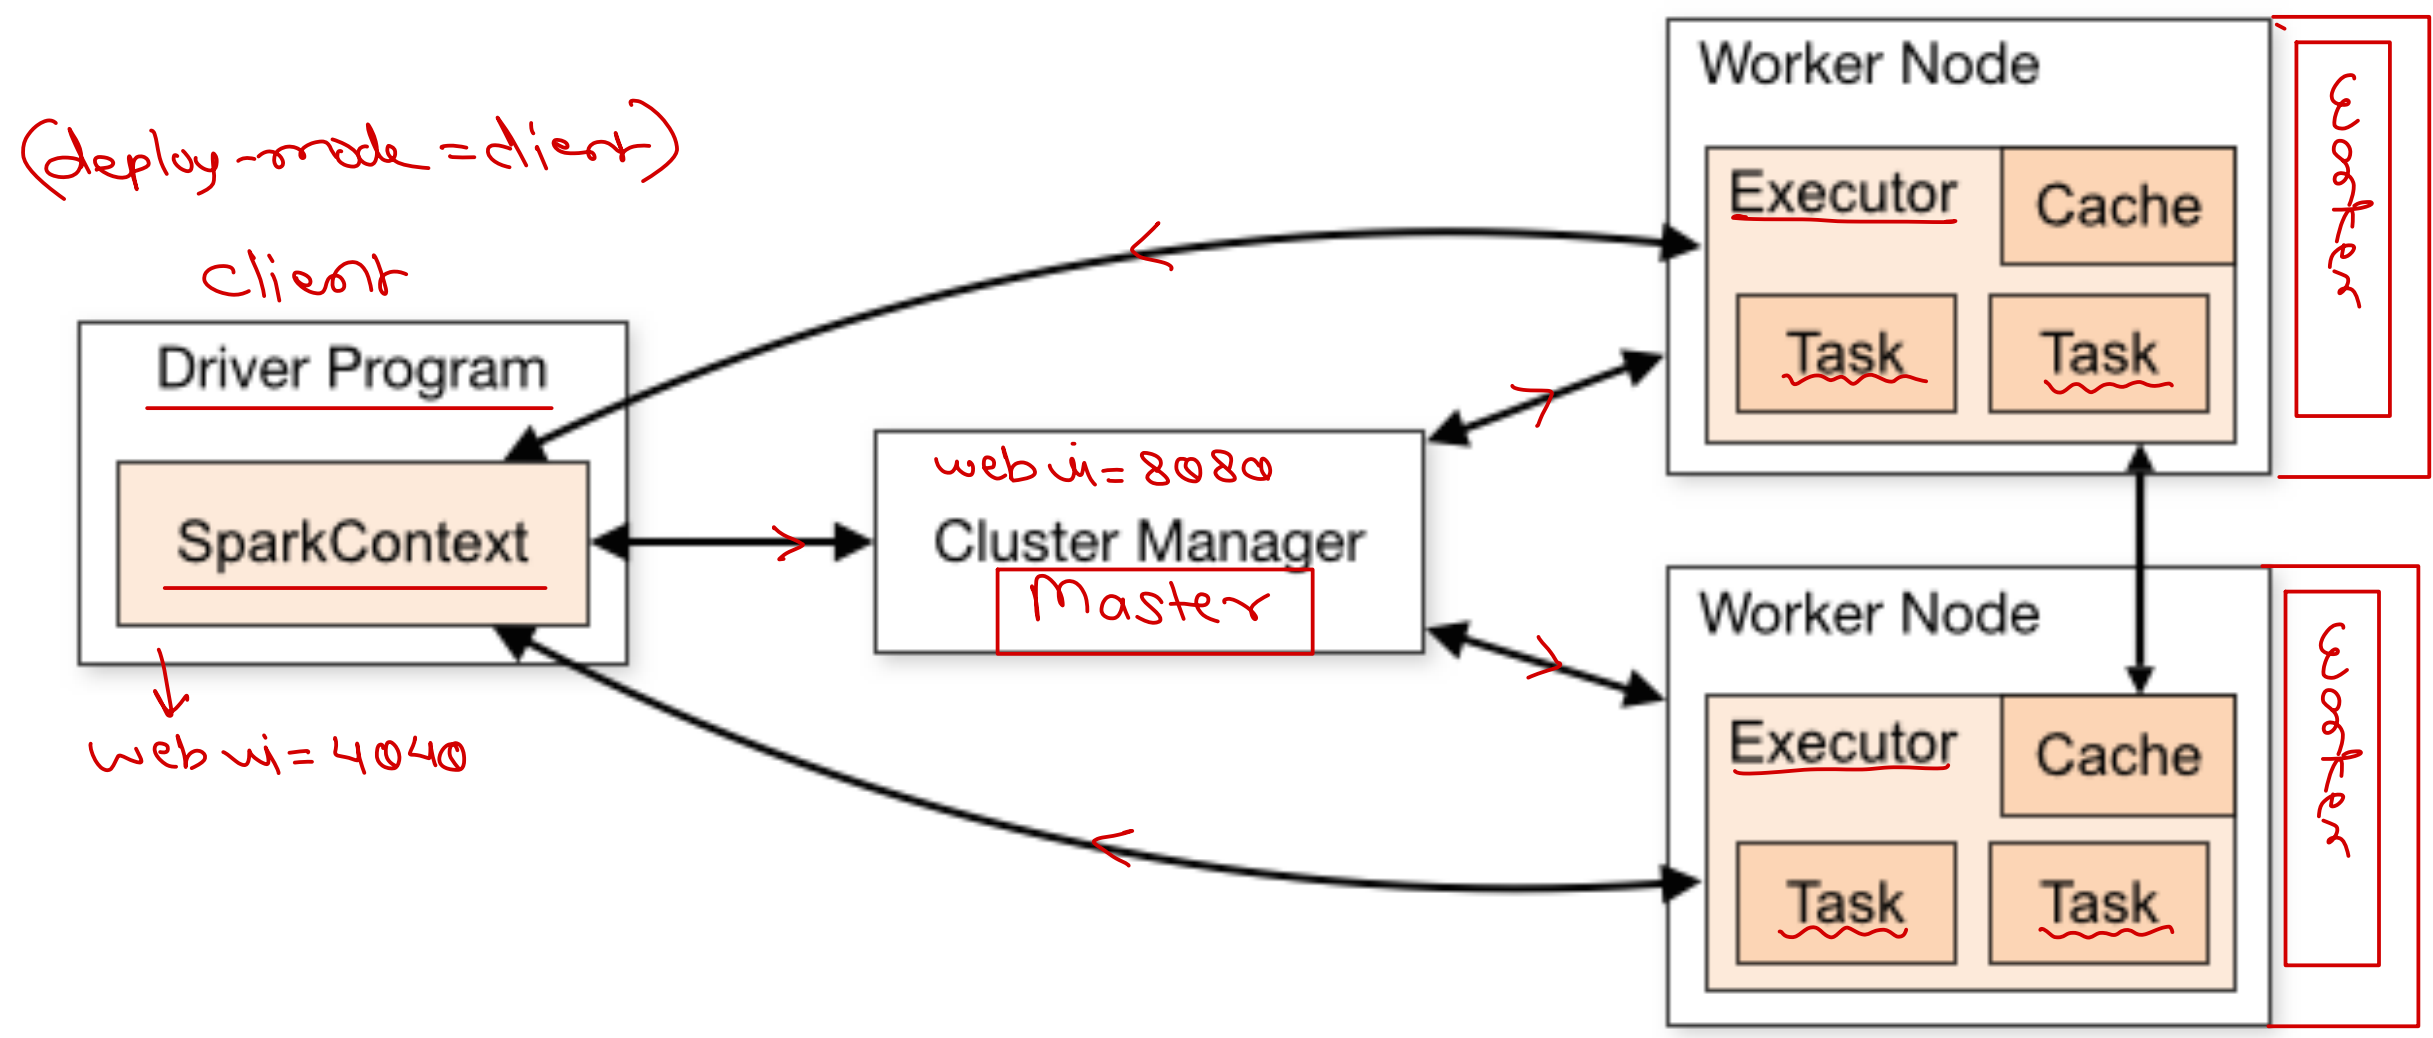
- terminal> start-master.sh
- terminal> ~~start-slaves.sh~~ start-workers.sh
- terminal> jps

- Using cluster:

- pyspark --master spark://localhost:7077 ✓
- spark-submit




Spark Cluster Manager - Spark Standalone mode



Spark Execution – Terminologies

- RDD: Resilient Distributed Dataset
 - Distributed across nodes -- Partition.
- DAG: Execution plan
 - Graph: Nodes=RDDs, Edges=Operations
- SparkContext: In-charge of execution
 - Maintain all info for execution/communication
 - Creates RDDs
- Application: Set of jobs
- Job: A DAG of execution
 - An action triggers job
- Stages: DAG is divided into stages
 - Stages are separated by shuffle operations (wide transformations)
- Task: Set of operation performed in a stage on a partition
 - A thread is created for each task.




Spark Execution – Terminologies

- Driver: Process that creates SparkContext
 - Create DAG
 - Submit DAG to the Master for execution.
 - Keep track of execution progress.
- Master: Cluster manager
 - It is a Java process. *web ui = 8080*
 - It manage and allocate resources.
 - Spark standalone mode: Master URL = spark://master:7077 *← IPC port*
- Worker: Each node in cluster
 - It is a Java process. *web ui = 8081^x*
 - It executes application in separate process -
- executor.
- Executor
 - It is a Java process.
 - Created by the worker for execution of *am*
application.
 - All tasks (threads) are created in executor.

→ created on client or worker node
depending --deploy-mode





Spark Multi Node Cluster (Standalone)

- On Master machine, conf/~~slaves~~^{workers} make entries of all slaves.
- In all machines, conf/spark-defaults.conf
 - spark.master spark://master:7077
- On master machine, set SPARK_LOCAL_IP & SPARK_MASTER_HOST to be set (in conf/spark-env.sh) to the IP address of network.
 - export SPARK_LOCAL_IP=master
 - export SPARK_MASTER_HOST=master
- On each slave machine, set SPARK_LOCAL_IP (in conf/spark-env.sh) to the IP address of network.
 - export SPARK_LOCAL_IP=~~vmx~~ worker ?
- from master
 - terminal> start-master.sh
 - terminal> start ~~slaves.sh~~ - workers.sh
- Using spark cluster
 - terminal> spark-submit --master spark://master:7077 --deploy-mode client app.py

- ✓ one master & three worker nodes
- ✓ in /etc/hosts → make entries of master & workers
- ✓ send ssh id of master to all workers.
- ✓ download & extract spark in home dir of same user (hduser).
- ✓ in .bashrc of all make \$SPARK_HOME & \$PATH entries.
- ✓ change hostname of all machines.





Thank you!

Nilesh Ghule <nilesh@sunbeaminfo.com>

