



Spark SQL

Sunbeam Infotech



Introduction

- Based on Spark structured API i.e. dataframes.
- Enable writing SQL queries on Spark dataframes as views/tables.
- Before Spark 2.x, SQLContext provides SQL functionality.
- Spark 2.x SparkSession encapsulate SparkContext. *+Sql Context*
- SparkContext use Hive metastore to maintain metadata.

abstraction = Catalog



Spark Tables

- Spark dataframes can be saved as table.
 - `df.saveAsTable("tablename")`
 - Table metadata is stored in metastore and data stored in spark warehouse directory.
- Spark tables can be partitioned by one or more column.
 - `df.write.partitionBy("col_name").saveAsTable("part_tablename")`
 - Partitions are sub-directories (directory name col=value) in which data is divided by column value.
- Spark tables can be bucketed by a column.
 - `df.write.bucketBy(numOfBuckets, colname).saveAsTable("buck_tablename")`
 - Buckets divide data into multiple data files by column value.
- Spark tables can be partitioned as well as bucketed.
 - `emp.write. partitionBy("col1").bucketBy(numOfBuckets, col2).saveAsTable("tablename")`
- Buckets are supported only as spark managed tables.



Spark Views

- View is abstraction on spark dataframes.
 - Created using df.createOrReplaceTempView("viewName")
 - createOrReplaceTempView()
 - Creates view if not available.
 - If available, replace with new view.
 - View treats dataframe as in memory table & create a view (like SQL view) to fire SQL queries on it.
 - The temporary view is in memory only, its info not stored in metastore. It is attached to current sparkSession.
 - df.createOrReplaceGlobalTempView("viewName") creates global view, which can be shared across multiple sessions.
- ndf = spark.sql("select from viewName")*



Spark SQL – setup

- Copy hive-site.xml into \$SPARK_HOME/conf
 - javax.jdo.option.ConnectionURL = spark/hive metastore path (derby/mysql)
 - javax.jdo.option.ConnectionDriverName = derby/mysql driver
 - javax.jdo.PersistenceManagerFactoryClass = persistence manager factory
 - hive.metastore.warehouse.dir/spark.sql.warehouse.dir = local/hdfs directory path
- Start spark master and slaves.
 - start-master.sh
 - start-slaves.sh
- Start spark thrift-server.
 - start-thriftserver.sh
- Start spark beeline.
 - beeline -u jdbc:hive2://localhost:10000 -n \$USER



Spark Hive Integration

Spark 3.3.x → hive 2.3.9.

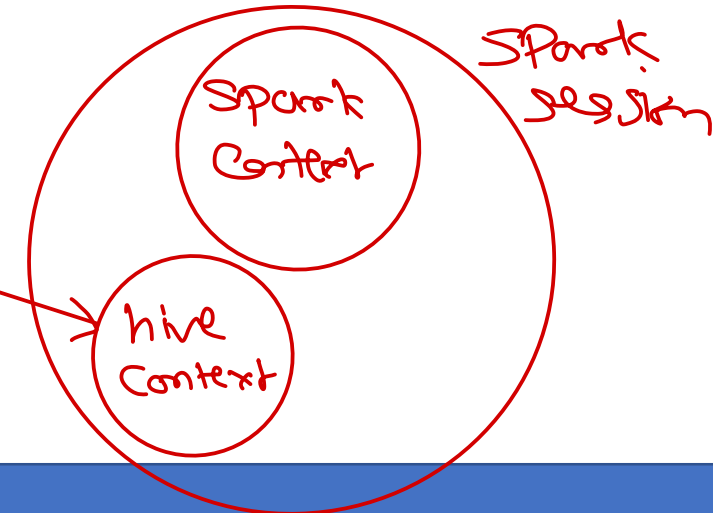
- Spark ^{2.3} metastore is compatible with Hive 1.2.1.
- Spark can access tables from Hive directly. However all dependencies of Hive are not shipped with Spark.
- To access Hive tables from spark application, Hive config should be associated with application and HiveContext should be activated.

```
spark = SparkSession.builder.appName("app")\  
    .config("javax.jdo.option.ConnectionURL", "jdbc:derby;;databaseName=/path/to/metastore")\  
    .config("javax.jdo.option.ConnectionDriverName", "org.apache.derby.jdbc.EmbeddedDriver")\  
    .config("hive.metastore.warehouse.dir", "/path/to/spark-warehouse")\  
    .enableHiveSupport().getOrCreate()
```

```
tables = spark.catalog.listTables()
```

```
books = spark.read.table("sbooks")
```

```
spark.stop();
```





Spark Streaming

Sunbeam Infotech



Batch processing vs Stream processing

- Processing finite set of data (data at rest).
- Incremental data load is managed by programmer.
- Cluster should be planned as per data size. High throughput.
- Job run once by batch.
- Processing live stream of data (data in motion).
- Data processing is managed by framework. *e.g. spark, flink, storm, ...*
- Less throughput.
- Job is running forever.



Stream processing

- Applications

- Notifications & Alerts: Shipping alert, Fire alert, ...
- Incremental ETL: Load live data from twitter/fb and process, ...
- Real time reporting: Live dashboard, ...
- Real time decisions: Customer management, ...
- Online ML: Training ML model with live data, fraud detection, ...

- Advantages

- Batch processing need to execute periodically (manually or scheduler).
- Processing with lower latency.
- Efficient handling of Incremental data.



Stream processing

- Challenges of stream processing

- Maintain large amount of state.
- Data throughput.
- Exactly once processing.
- Process out-of-order data.
- Low latency processing.
- Load imbalance.
- Join with external data.
- Producing output.

- Design considerations


- Record at a time vs Declarative APIs
- Event time vs Processing time
- Continuous processing vs Micro-batch processing

Record processing

- ① Exactly once
- ② at least once
- ③ at most once

R1 - 12:28:15
R2 - 12:28:20
R3 - 12:28:25
R4 - 12:28:30

event time




R1 - 12:28:17
R2 - 12:28:29
R3 - 12:28:28
R4 - 12:28:32

processing time

low latencies
but difficult handling } storm

easier handling
but higher latencies } spark



Spark Streaming

- Spark is originally designed ^{for} by micro-batch processing.
- Spark Streaming APIs
 - Spark DStream ✓
 - Spark Structured Streaming ✓



Spark DStream

- Micro-batches of RDDs (Small RDDs).
- Developed in 2012. Most popular Streaming framework in 2016.
- RDD based programming.
- Limitations
 - Based on RDDs (in JVM). Not efficient in Python.
 - No support for event time processing.
 - Only micro-batch processing.
- Examples
 - Twitter stream
 - Socket stream processing



Spark Structured Streaming

- Developed in 2016.
- Stable in Spark 2.2.
- Spark Structured Streaming is based on dataframes.
- Works seamlessly with other Spark APIs i.e. Spark SQL & Spark ML.
- Advantages
 - Optimized (Catalyst engine)
 - Event time processing is supported
 - Support for continuous processing
 - Same query/code works for batch processing & stream processing
 - Exactly once processing mode is available
 - Fault tolerance



Spark Structured Streaming

- Spark Structured Streaming consider dataframe to be unbounded (infinitely growing).
- Transformations & Actions
 - Transformations are same as spark dataframe. Few transformations are not yet implemented.
 - Action is starting the stream & print results.
- Input sources
 - socket, rate, files, flume, kinesis, kafka
- Output sinks
 - console, memory, files, flume, kafka, foreach



Spark Structured Streaming

- Output modes
 - Every mode is not supported for every type of query.
 - append: output result of current micro-batch is available. (not supported for aggregate operations).
 - complete: complete result including prev result & current micro-batch result is available.
 - update: only results modified in current micro-batch are available.
- Triggers
 - By default, micro-batches are processed one after another.
 - Trigger can specify time duration after which each batch is to be processed.
- Event time processing
 - Time at which event is generated at source, is "event time".
 - Can process out-of-order data.
 - Watermark feature is used define for how much time data should be considered (how much time should be wait before processing data).





Thank you!

Nilesh Ghule <nilesh@sunbeaminfo.com>

