



Sunbeam Infotech



Introduction

Free & Open source

- A platform created by community to programmatically author, schedule and monitor workflow. → set of tasks to be done is a specific sequence. Python → horizontal scaling → web ui (port 8080)
- Airflow is used to develop, orchestrate and monitor complex ETL pipelines.
- Airflow is completely developed in Python, so it can work with any of the Python library. → kafka-python, sklearn, ..

→ github - 29K stars

- Initiated by Airbnb in 2014 and open-sourced under Apache in 2016.
- Apache top level project in 2019.
- Used by more than 250 companies world-wide - Amazon, Citi, JPMorgan, Salesforce, Drillinginfo, Unitedhealth, ...
- Managed services in GCP and AWS cloud.

AWS managed services

① EMR

⑤ S3

⑨ Beanstalk

② Redshift

⑥ RDS

⑩ Airflow

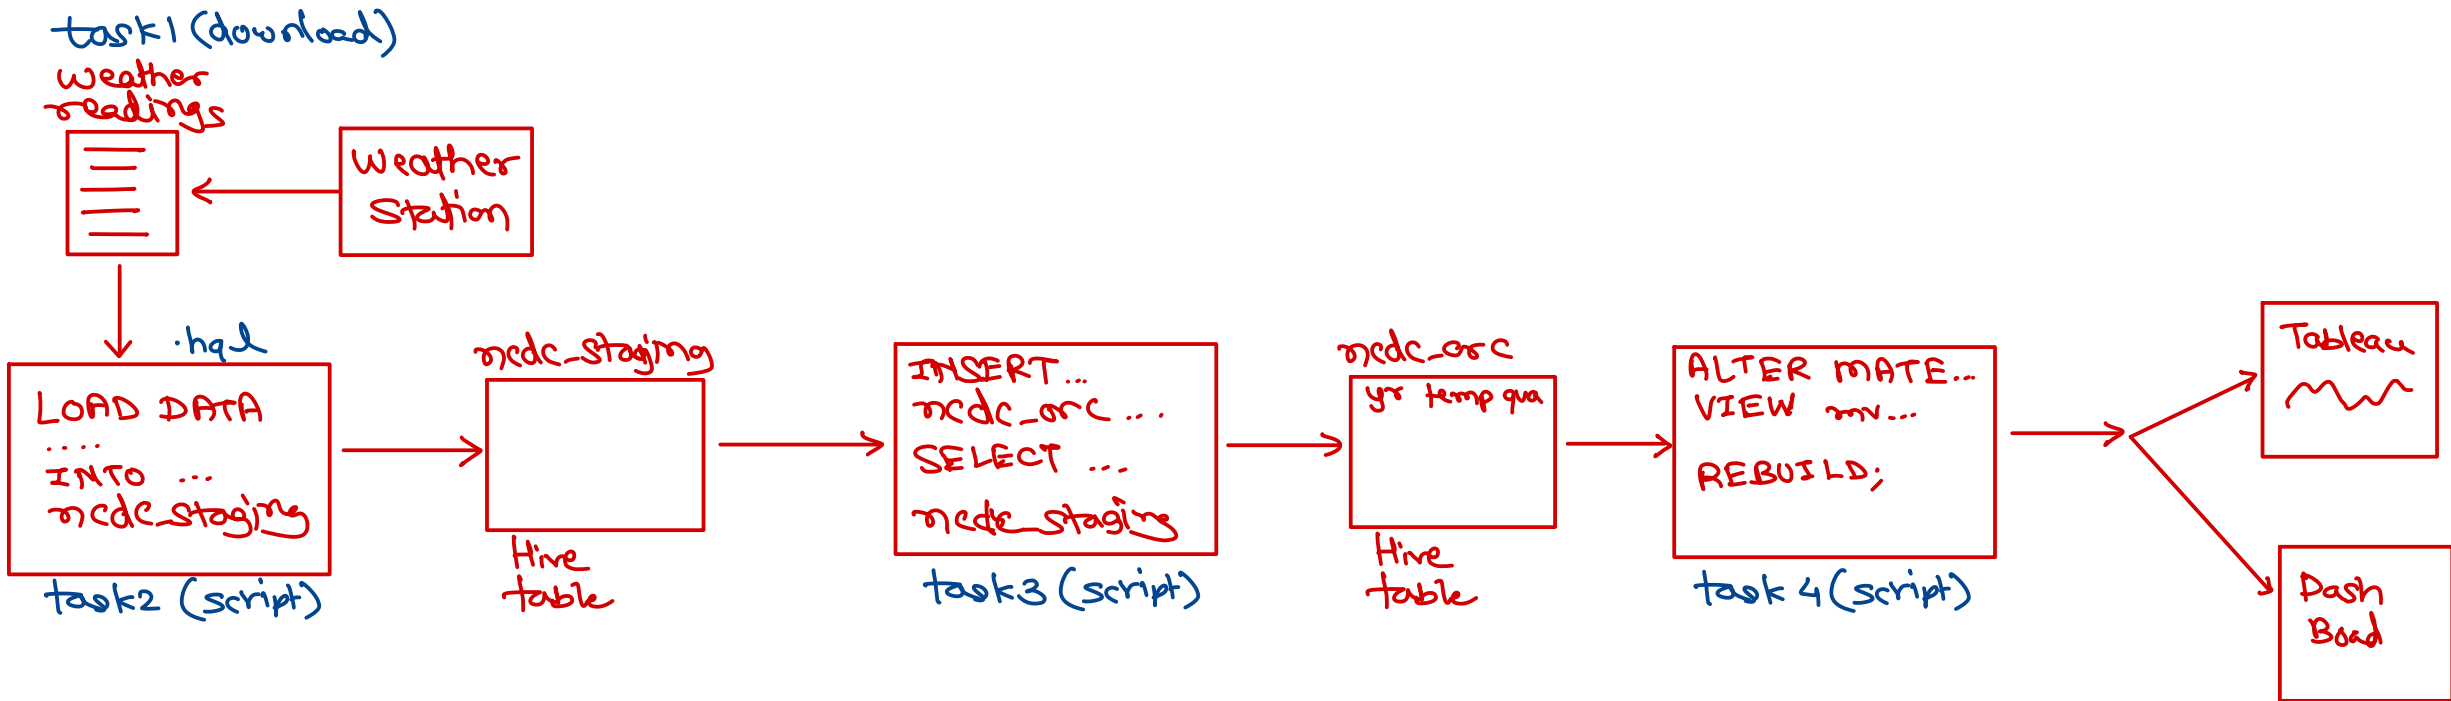
③ EKS

⑦ CloudWatch

④ Glue

⑧ Kinesis

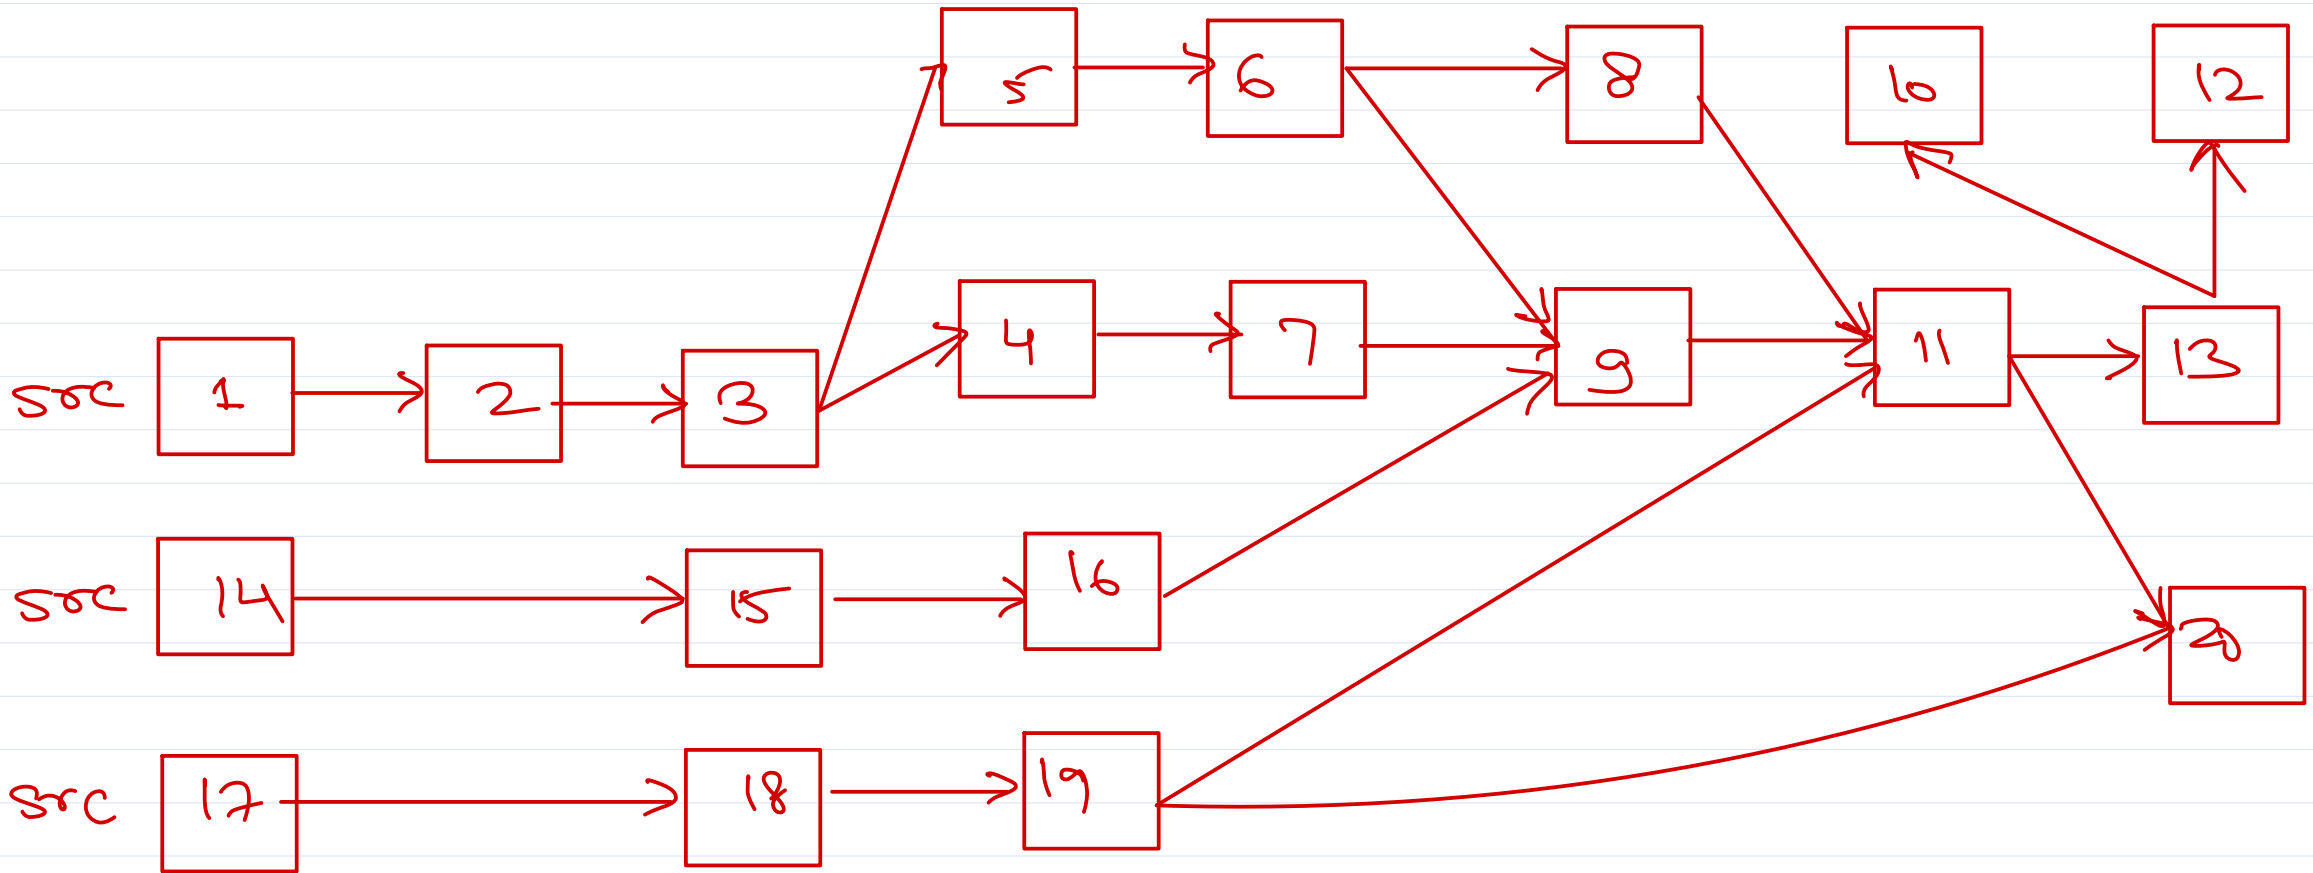




Traditional ETL

- Extract, Transform and Load
- Common schedulers: cron, oozie (hadoop), luigi (spotify)
- BigData has complex pipelines.
- Cron limitations
 - Error handling
 - Not maintainable (runs on a machine) *no scaling*
 - Execution dependency (delayed tasks)
 - Transparency (No centralized log)
 - Task tracking (No centralized monitoring)
 - Handle historical data
- Oozie
 - Error handling (n retries)
 - Execution dependency
 - Better tracking & monitoring
 - XML based jobs - difficult to read/understand





Airflow Terminologies

- DAG

- Directed Acyclic Graph
- Vertices are tasks and Edges are dependencies
- Can be parallelized.

$t_1 \rightarrow t_2 \rightarrow t_3 \rightarrow t_4$

$t_1 \rightarrow t_2 \rightarrow \begin{cases} t_3 \rightarrow t_4 \\ t_5 \rightarrow t_6 \end{cases} \rightarrow t_7 \rightarrow t_8$

- Operators

- Single dedicated task in workflow/DAG.
- Example: Bash command, Python Function, Database Operation, Email send, etc.

- Task


- Instantiated operator and assigned to some worker.
- Failed task can be retried.

- Workflow

- Sequence of task arranged in control dependencies.
- Workflow and DAG words are interchangeable.

business
terminology

programming
terminology



Airflow Workflow

- Airflow workflows are written as Python code in form of DAG.
 - DAG includes multiple tasks in form of operators.
 - Operators are connected in complex sequence.
- DAG implemented in Python, but can execute variety of tasks.
 - Programs written in any language.
 - Analytical/ML task → Python Operator - sklearn
 - ETL task
 - Big Data components/programs → Hive Operator, ...
- DAGs are more dynamic, manageable, testable and collaborative.
- Airflow is job scheduler that execute the tasks on defined time and/or followed by its dependencies.



Airflow Advantages

- Dynamic DAG
 - Implemented in code - flexible
 - Can be configured: parallelism, params, templates → *ninja {{ ds }}*
- Extensible
 - Plenty of operators/executors
 - Shell script, Big data task, OS command, etc.
 - Custom task/plugins
- Scalable
 - Modular architecture
 - Can handle any number of DAGs (with multi-node cluster)
- Configurable
 - airflow.cfg - admin settings
 - Centralized configurations
- Monitoring
 - Elegant Web UI 8080
 - Handle failures (n retries)
 - Email alerts
- Open Source
 - Active community
 - New plugins



Airflow Applications

- Task scheduler
- ETL pipelines
- Periodic backups
- Automate DevOps operations
- ML jobs
- Recommendation engines
- ...

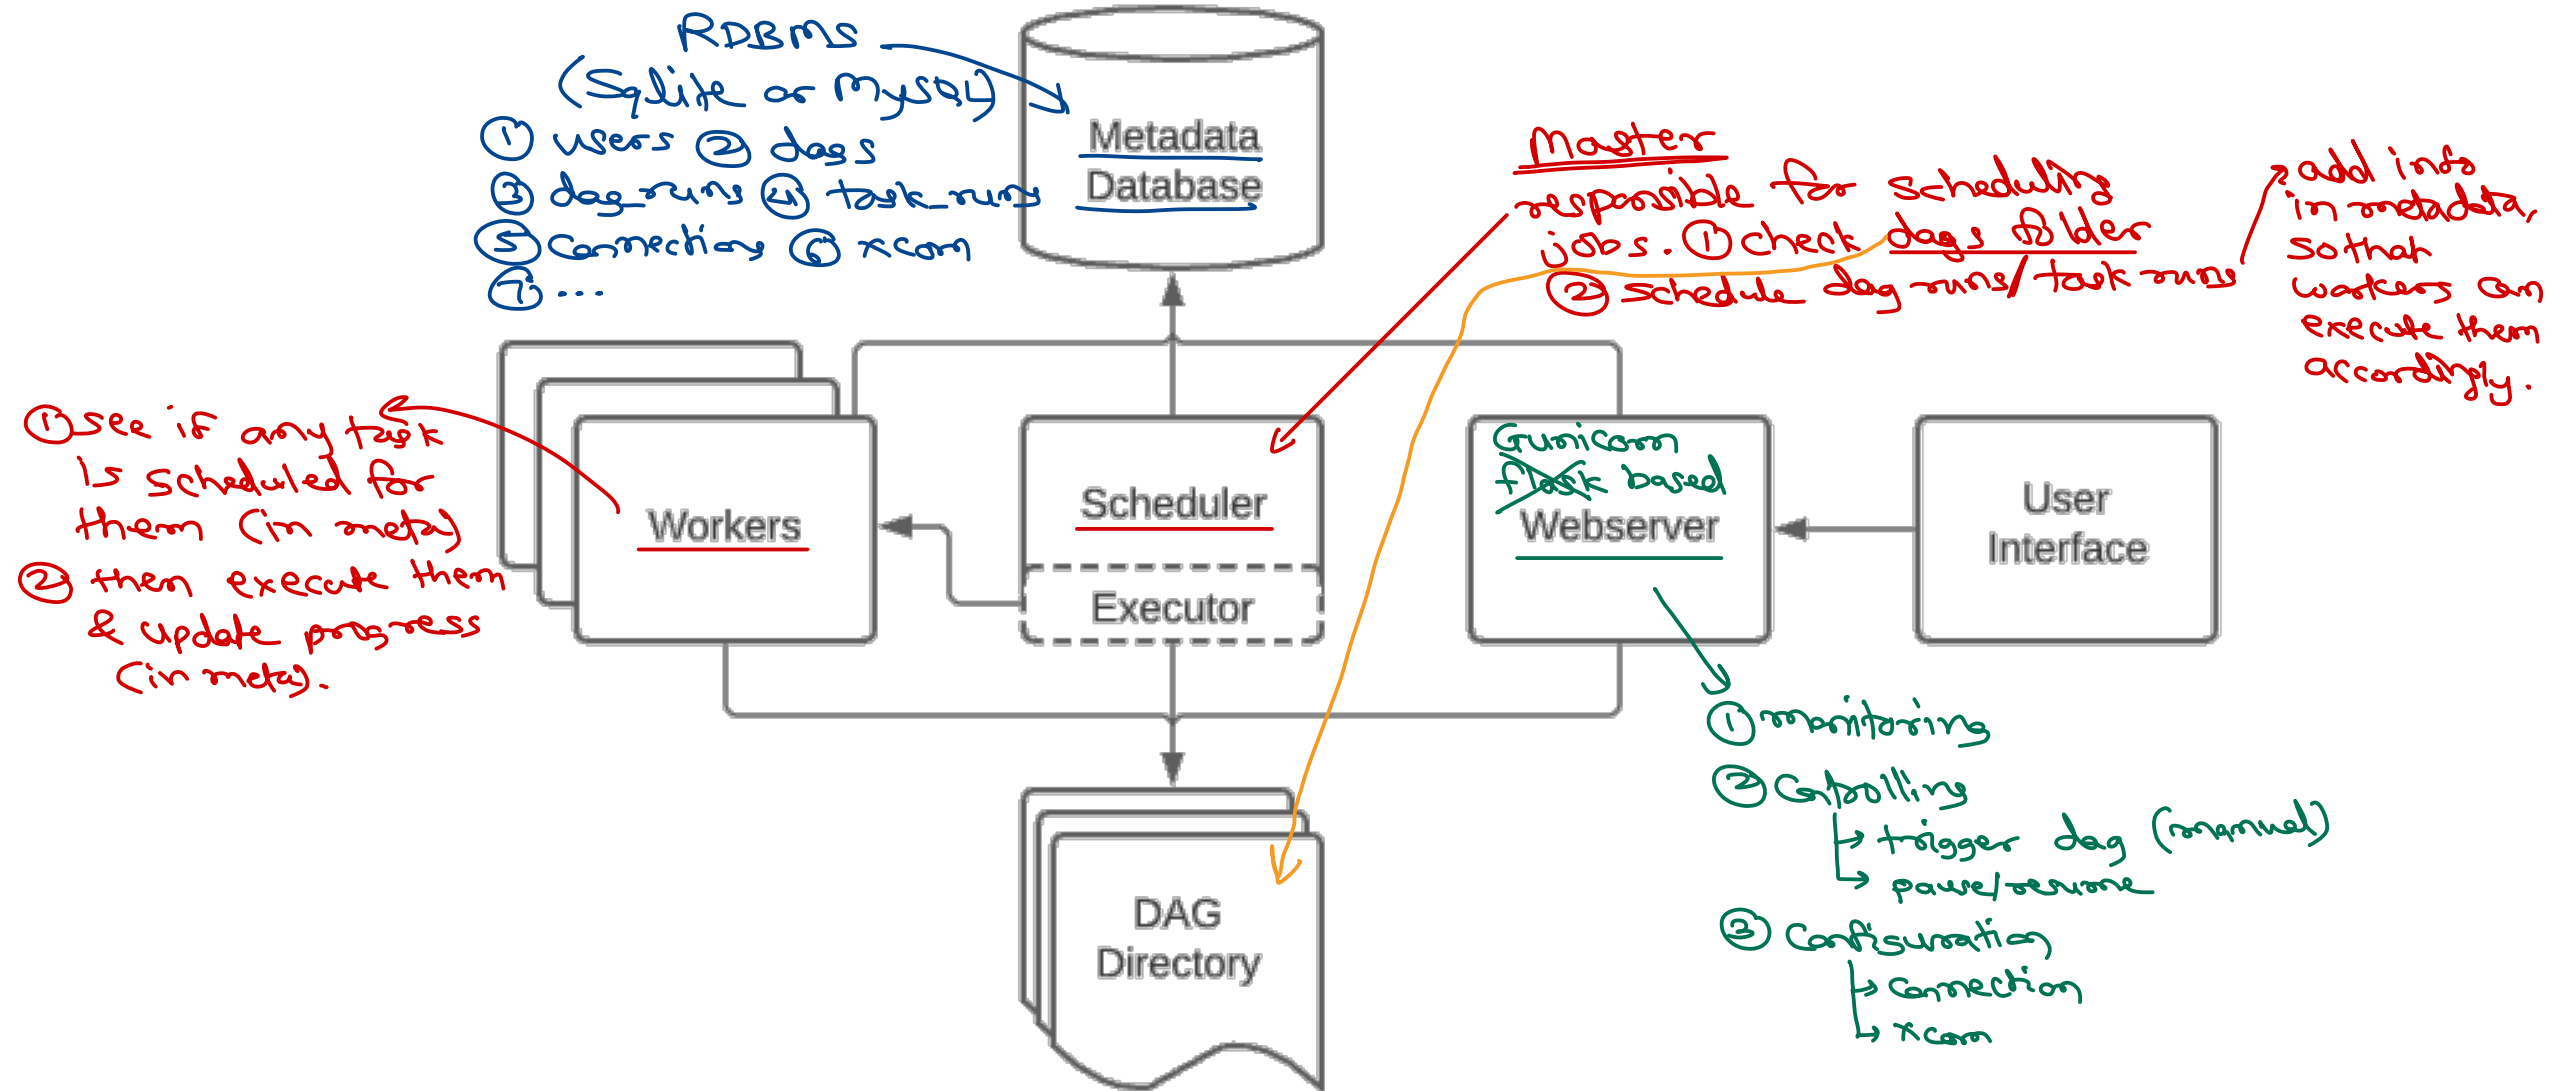


Airflow Installation

- Airflow is installed as Python package – refer docs.
- Airflow directory/files *HOME → /home/username/airflow*
 - config
 - *✓* airflow.cfg
 - dags folder → *\$AIRFLOW_HOME/dags*
 - timezone → *default = UTC*
 - executor type → *default = Sequential*
 - parallelism → *default = 32*
 - *dags → directory* ← *(points from dags folder in config)*
 - dag definition (python files)
 - script
 - *✓* entrypoint.sh (with docker image)
 - to initiate airflow - web-server & scheduler.
 - *✓* environment variables
 - metastore
 - airflow.db *↖ default = sqlite*
 - using sqlite (default) database for airflow metastore.



Airflow Architecture



Airflow Architecture

- Metadata

- RDBMS that stores historical and current DAGs and tasks details.
- Also maintains information used resources.
- Default RDBMS used is SQLite, but it is single-user database.
- In production MySQL or Posgre-SQL is preferred option.

- Scheduler

- Instructs and trigger task execution on worker node.
- Implemented into Python.
- Reads DAG file, task config and schedules the task on worker nodes in sequence.
- It monitor tasks state from metadata and handle the task failure (as per config).

- Web Server

- Flask-based UI for Airflow to monitor DAG.
- Communicate with metadata to fetch the task state.
- Task state is rendered in various formats including DAG, graphs, time/numbers, etc.



Airflow Architecture

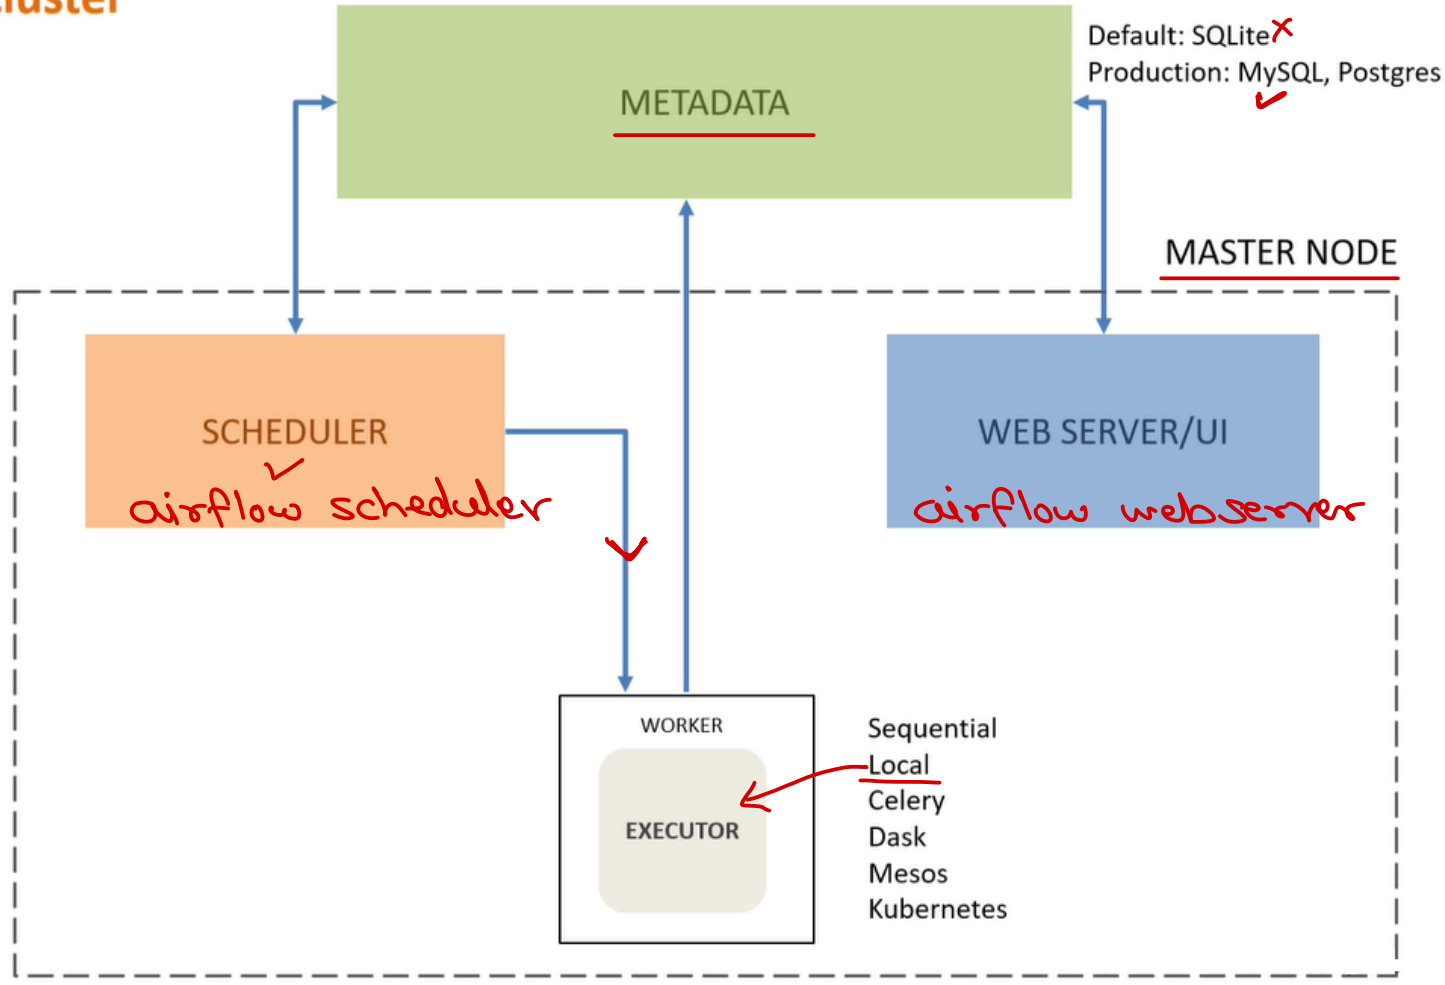
- Executor

- Carry the task scheduled by the scheduler.
- Runs on Worker node.
- Various types:
 - Sequential (default) - For testing
 - Local - For single node cluster - Use python process. ✓
 - Celery - Recommended for multi-node cluster. ✓
 - Job queue written in Python.
 - Used in scalable distributed system.
 - Dask
 - Mesos
 - Kubernetes ✓



Airflow – Single node installation

Single Node cluster

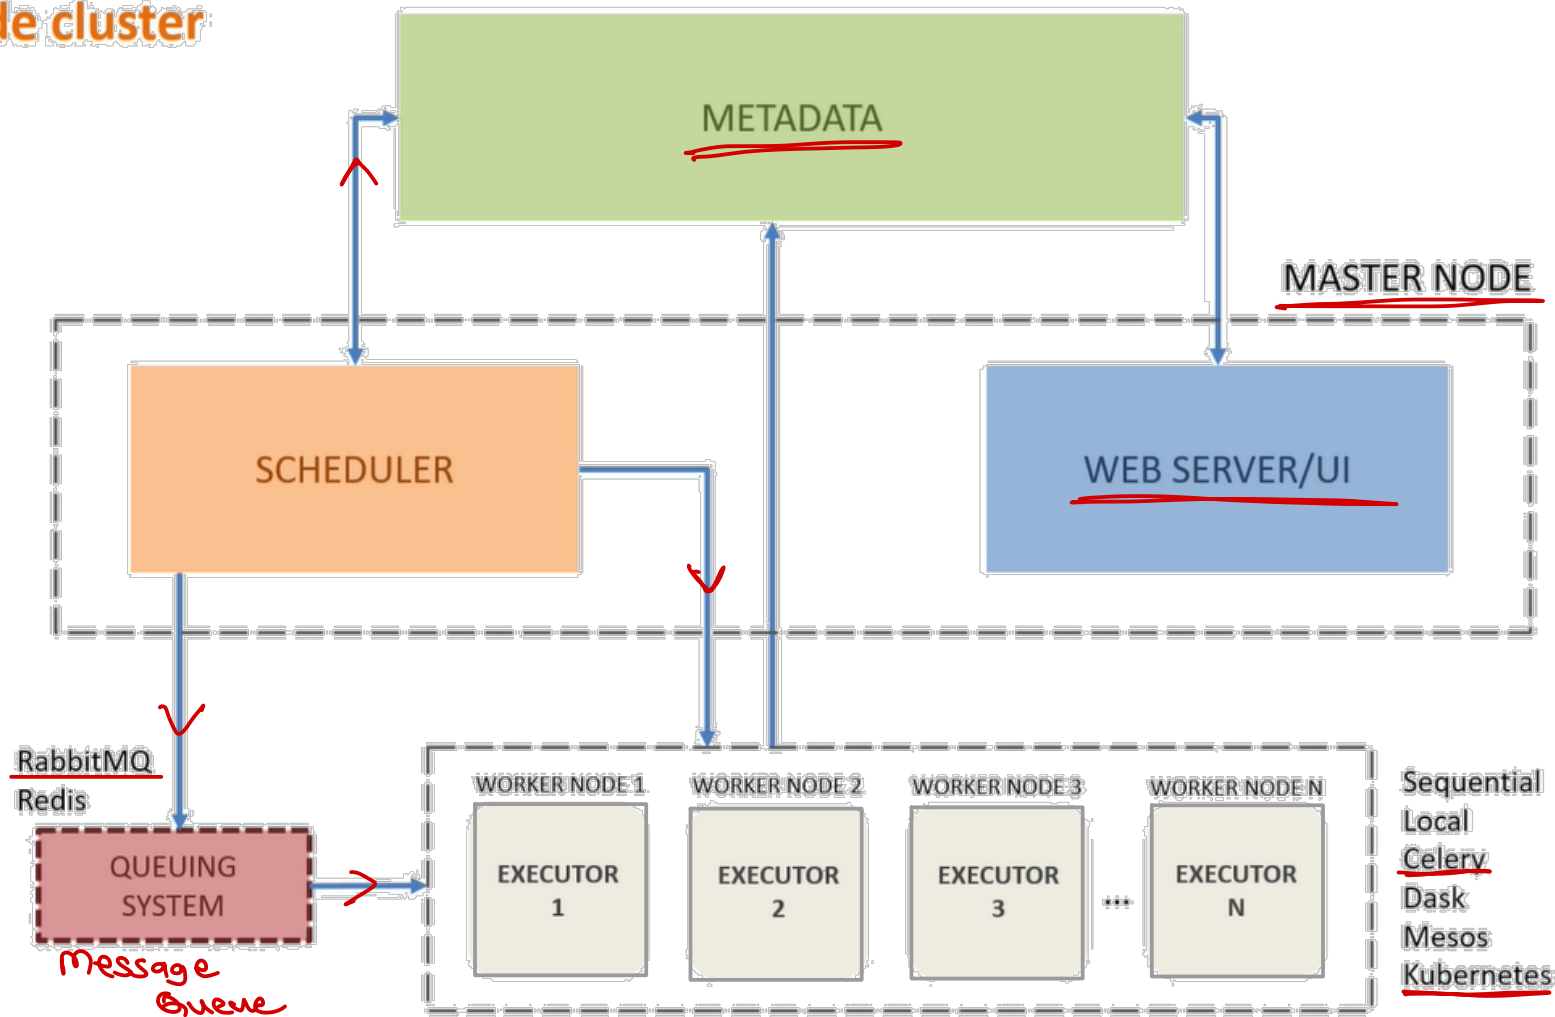


- Single worker node.
- Not scalable (max all resources available on the system).
- Quite fast for limited number of DAGs and data size.
- Use local executor.
- Direct communication between scheduler and executor.



Airflow – Multi node installation

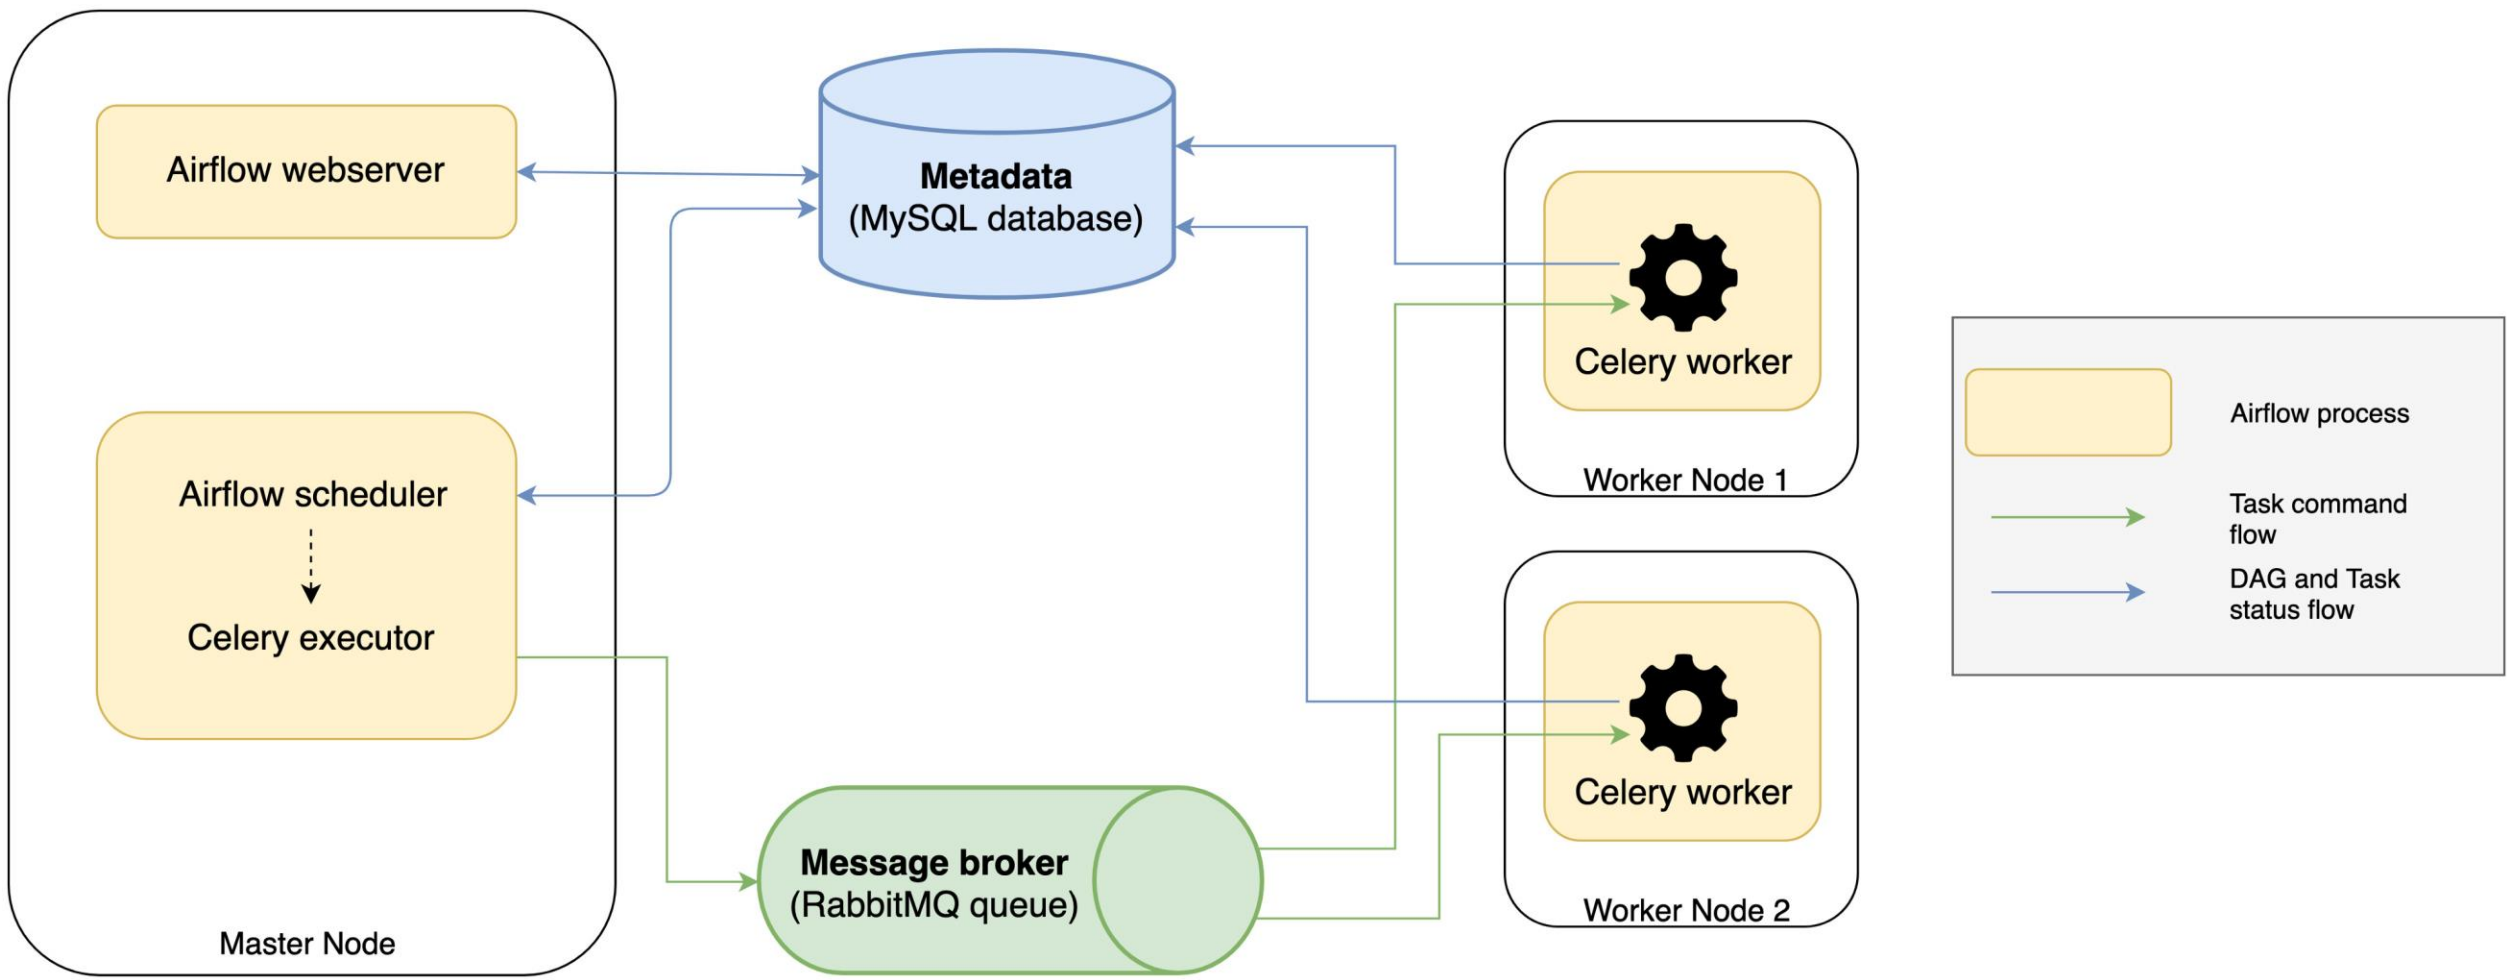
Multi Node cluster



- Master node runs scheduler and web-server.
- Multiple worker nodes - running executors.
- Highly scalable for huge data and many nodes.
- Recommended executor is Celery.
- Scheduler and workers communicate with external queue system like Rabbit MQ / Redis.



Airflow – Multi node installation



Airflow Working

- Scheduler periodically pings for DAG folder and communicate with metastore.
- If any DAG is available for execution, scheduler starts a DAG run for it. DAG run is an object representing an instantiation of DAG.
- Scheduler update DAG state as "Running".
- For each task, task object is instantiated. The task state is updated as "Scheduled".
- Scheduler assigns priority to each task in DAG (as per config) and push them into queuing system. The task state is updated to "Queued".
- Worker pull the task from queue, set its state as "Running" and start executing it.
- Upon completion of each task, task state is updated as succeed or failed.
- When all tasks in DAG run are executed, status of DAG run is updated as succeed or failed.
- Airflow web-server periodically get the data from meta-store and render it for users.
- If any new DAG is found in DAGs folder, scheduler begin its execution (as above).



Airflow Operators

- DAG only describe how to run the workflow, but do not perform actual computation.
- Operator describes single task in workflow.
- Operator characteristics
 - Atomic (usually) - standalone (not sharing resources with other operators)
 - Idempotent - produce same result for each run.
 - Operator in execution (instantiated) is referred as task.
 - Inherited from airflow's BaseOperator class.
 - Operators can communicate using XCom.
- Operator categories
 - Sensor operators - Wait for certain criteria
 - HdfsSensor, FileSensor, ...
 - Transfer operators - Transfer the data
 - MySqlToHiveOperator, ...
 - Action operators - Do specified task
 - BashOperator, PythonOperator, HiveOperator, MySqlOperator, EmailOperator, ...





Thank you!

Nilesh Ghule <nilesh@sunbeaminfo.com>

